

# STAD29: Statistics for the Life and Social Sciences

Lecture notes

## Section 1

# Principal components

# Principal Components

- Have measurements on (possibly large) number of variables on some individuals.
- Question: can we describe data using fewer variables (because original variables correlated in some way)?
- Look for direction (linear combination of original variables) in which values *most spread out*. This is *first principal component*.
- Second principal component then direction uncorrelated with this in which values then most spread out. And so on.

# Principal components

- See whether small number of principal components captures most of variation in data.
- Might try to interpret principal components.
- If 2 components good, can make plot of data.
- (Like discriminant analysis, but no groups.)
- “What are important ways that these data vary?”

# Packages

You might not have installed the first of these. See over for instructions.

```
library(ggbiplot) # see over  
library(tidyverse)  
library(ggrepel)
```

# Installing ggbiplot

- ggbiplot not on CRAN, so usual `install.packages` will not work. This is same procedure you used for `smmr` in C32:
- Install package `devtools` first (once):

```
install.packages("devtools")
```

- Then install `ggbiplot` (once):

```
library(devtools)  
install_github("vqv/ggbiplot")
```

## Small example: 2 test scores for 8 people xxx


```
my_url <- "http://www.utsc.utoronto.ca/~butler/d29/test12.txt"
test12 <- read_table2(my_url)
test12
```

```
## # A tibble: 8 x 3
##   first second id
##   <dbl>  <dbl> <chr>
## 1      2      9 A
## 2     16     40 B
## 3      8     17 C
## 4     18     43 D
## 5     10     25 E
## 6      4     10 F
## 7     10     27 G
## 8     12     30 H
```

```
g <- ggplot(test12, aes(x = first, y = second, label = id)) +
  geom_point() + geom_text_repel()
```

# The plot

```
g + geom_smooth(method = "lm", se = F)
```





# Principal component analysis

- Grab just the numeric columns:

```
test12 %>% select_if(is.numeric) -> test12_numbers
```

- Strongly correlated, so data nearly 1-dimensional:

```
cor(test12_numbers)
```

```
##           first    second
## first  1.000000  0.989078
## second 0.989078  1.000000
```

# Finding principal components

- Make a score summarizing this one dimension. Like this:

```
test12.pc <- princomp(test12_numbers, cor = T)
summary(test12.pc)
```

```
## Importance of components:
```

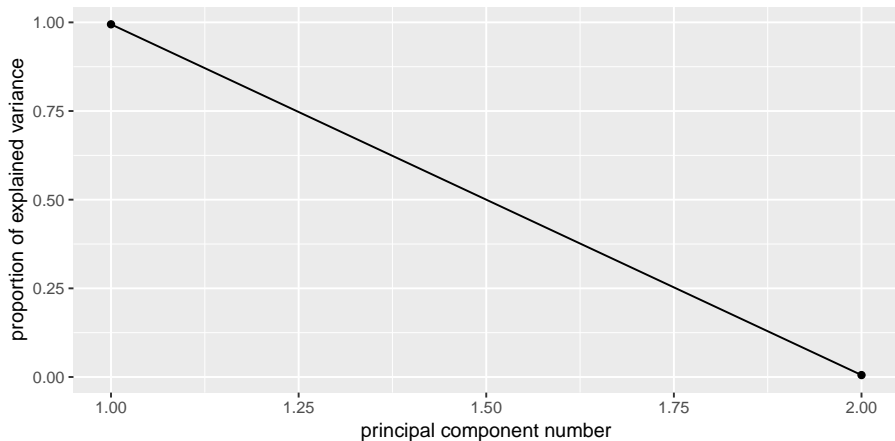
##	Comp.1	Comp.2
## Standard deviation	1.410347	0.104508582
## Proportion of Variance	0.994539	0.005461022
## Cumulative Proportion	0.994539	1.000000000

# Comments

- “Standard deviation” shows relative importance of components (as for LDs in discriminant analysis)
- Here, first one explains almost all (99.4%) of variability.
- That is, look only at first component and ignore second.
- $\text{cor=T}$  standardizes all variables first. Usually wanted, because variables measured on different scales. (Only omit if variables measured on same scale and expect similar variability.)

# Scree plot

```
ggscreeplot(test12.pc)
```



# xxx Component loadings

explain how each principal component depends on (standardized) original variables (test scores):

```
test12.pc$loadings
```

```
##
## Loadings:
##          Comp.1 Comp.2
## first    0.707  0.707
## second   0.707 -0.707
##
##          Comp.1 Comp.2
## SS loadings      1.0    1.0
## Proportion Var   0.5    0.5
## Cumulative Var   0.5    1.0
```

First component basically sum of (standardized) test scores. That is, person tends to score similarly on two tests, and a composite score would summarize performance.

## xxx Component scores


```
d <- data.frame(test12, test12.pc$scores)
d
```

```
##   first second id      Comp.1      Comp.2
## 1      2      9  A -2.071819003 -0.146981782
## 2     16     40  B  1.719862811 -0.055762223
## 3      8     17  C -0.762289708  0.207589512
## 4     18     43  D  2.176267535  0.042533250
## 5     10     25  E -0.007460609  0.007460609
## 6      4     10  F -1.734784030  0.070683441
## 7     10     27  G  0.111909141 -0.111909141
## 8     12     30  H  0.568313864 -0.013613668
```

- Person A is a low scorer, high positive comp.1 score.
- Person D is high scorer, high negative comp.1 score.
- Person E average scorer, near-zero comp.1 score.
- comp.2 says basically nothing.

# Plot of scores

```
ggplot(d, aes(x = Comp.1, y = Comp.2, label = id)) +  
  geom_point() + geom_text_repel()
```




# Comments

- Vertical scale exaggerates importance of comp.2.
- Fix up to get axes on same scale:

```
g <- ggplot(d, aes(x = Comp.1, y = Comp.2, label = id)) +  
  geom_point() + geom_text_repel() +  
  coord_fixed()
```

- Shows how exam scores really spread out along one dimension:

g






# The biplot

- Plotting variables and individuals on one plot.
- Shows how components and original variables related.
- Shows how individuals score on each component, and therefore suggests how they score on each variable.
- Add `labels` option to identify individuals:

```
g <- ggbiplot(test12.pc, labels = test12$id)
```

# The biplot



# Comments

- Variables point almost same direction (left). Thus very negative value on comp.1 goes with high scores on both tests, and test scores highly correlated.
- Position of individuals on plot according to scores on principal components, implies values on original variables. Eg.:
- D very negative on comp.1, high scorer on both tests.
- A and F very positive on comp.1, poor scorers on both tests.
- C positive on comp.2, high score on first test relative to second.
- A negative on comp.2, high score on second test relative to first.

# xxx Track running data

Track running records (1984) for distances 100m to marathon, arranged by country. Countries labelled by (mostly) Internet domain names (ISO 2-letter codes):

```
my_url <- "http://www.utsc.utoronto.ca/~butler/d29/men_track_field.txt"
track <- read_table(my_url)
track %>% sample_n(8)
```

```
## # A tibble: 8 x 9
##   m100  m200  m400  m800 m1500 m5000 m10000 marathon country
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl> <chr>
## 1  10.8  21.9   49    2.02  4.24  16.3   34.7    162. ws
## 2  10.3  20.7  45.0   1.73  3.6   13.2   27.4    130. be
## 3  10.4  20.6  45.6   1.76  3.58  13.4   28.2    134. cz
## 4  10.3  20.8  45.9   1.79  3.64  13.4   27.7    129. jp
## 5  10.6  21.5  47.8   1.84  3.92  14.7   30.8    149. id
## 6  10.4  20.8  46.8   1.81  3.7   14.0   29.4    138. ar
## 7  10.3  20.1  44.8   1.74  3.57  13.3   27.7    128. au
## 8  10.2  20.6  45.6   1.77  3.61  13.3   27.9    131. se
```

# xxx Country names

Also read in a table to look country names up in later:

```
my_url <- "http://www.uts.utoronto.ca/~butler/d29/isocodes.csv"
iso <- read_csv(my_url)
iso
```

```
## # A tibble: 251 x 4
##   Country      IS02 IS03   M49
##   <chr>      <chr> <chr> <dbl>
## 1 <NA>      <NA> <NA>   NA
## 2 Afghanistan af    afg     4
## 3 Aland Islands ax    ala   248
## 4 Albania    al    alb     8
## 5 Algeria    dz    dza    12
## 6 American Samoa as    asm    16
## 7 Andorra    ad    and    20
## 8 Angola     ao    ago    24
## 9 Anguilla   ai    aia   660
## 10 Antarctica aq    ata    10
## # ... with 241 more rows
```

# Data and aims

- Times in seconds 100m–400m, in minutes for rest (800m up).
- This taken care of by standardization.
- 8 variables; can we summarize by fewer and gain some insight?
- In particular, if 2 components tell most of story, what do we see in a plot?


## xxx Fit and examine principal components

```
track_num <- track %>% select_if(is.numeric)
track.pc <- princomp(track_num, cor = T)
summary(track.pc)
```

```
## Importance of components:
##
##              Comp.1      Comp.2
## Standard deviation    2.5733531 0.9368128
## Proportion of Variance 0.8277683 0.1097023
## Cumulative Proportion 0.8277683 0.9374706
##
##              Comp.3      Comp.4
## Standard deviation    0.39915052 0.35220645
## Proportion of Variance 0.01991514 0.01550617
## Cumulative Proportion 0.95738570 0.97289187
##
##              Comp.5      Comp.6
## Standard deviation    0.282630981 0.260701267
## Proportion of Variance 0.009985034 0.008495644
## Cumulative Proportion 0.982876903 0.991372547
##
##              Comp.7      Comp.8
## Standard deviation    0.215451919 0.150333291
## Proportion of Variance 0.005802441 0.002825012
```

# Scree plot

```
ggscreeplot(track.pc)
```





# How many components?

- As for discriminant analysis, look for “elbow” in scree plot.
- See one here at 3 components; everything 3 and beyond is “scree”.
- So take 2 components.
- Note difference from discriminant analysis: want “large” rather than “small”, so go 1 step left of elbow.
- Another criterion: any component with eigenvalue bigger than about 1 is worth including. 2nd one here has eigenvalue just less than 1.
- Refer back to summary: cumulative proportion of variance explained for 2 components is 93.7%, pleasantly high. 2 components tell almost whole story.

# xxx How do components depend on original variables?

Loadings:

```
track.pc$loadings
```

```
##
## Loadings:
##      Comp.1 Comp.2 Comp.3 Comp.4 Comp.5 Comp.6 Comp.7 Comp.8
## m100      0.318  0.567  0.332  0.128  0.263  0.594  0.136  0.106
## m200      0.337  0.462  0.361 -0.259 -0.154 -0.656 -0.113
## m400      0.356  0.248 -0.560  0.652 -0.218 -0.157
## m800      0.369          -0.532 -0.480  0.540          -0.238
## m1500     0.373 -0.140 -0.153 -0.405 -0.488  0.158  0.610  0.139
## m5000     0.364 -0.312  0.190          -0.254  0.141 -0.591  0.547
## m10000    0.367 -0.307  0.182          -0.133  0.219 -0.177 -0.797
## marathon  0.342 -0.439  0.263  0.300  0.498 -0.315  0.399  0.158
##
##      Comp.1 Comp.2 Comp.3 Comp.4 Comp.5 Comp.6 Comp.7
## SS loadings      1.000  1.000  1.000  1.000  1.000  1.000  1.000
## Proportion Var   0.125  0.125  0.125  0.125  0.125  0.125  0.125
## Cumulative Var   0.125  0.250  0.375  0.500  0.625  0.750  0.875
""
```

## xxx Comments

- comp.1 loads about equally (has equal weight) on times over all distances.
- comp.2 has large positive loading for short distances, large negative for long ones.
- comp.3: large negative for middle distance, large positive especially for short distances.
- Country overall good at running will have lower than average record times at all distances, so comp.1 *small*. Conversely, for countries bad at running, comp.1 very positive.
- Countries relatively better at sprinting (low times) will be *negative* on comp.2; countries relatively better at distance running *positive* on comp.2.

## xxx Commands for plots

- Principal component scores (first two). Also need country names.

```
d <- data.frame(track.pc$scores,
  country = track$country
)
names(d)
```

```
## [1] "Comp.1" "Comp.2" "Comp.3" "Comp.4" "Comp.5" "Comp.6"
## [7] "Comp.7" "Comp.8" "country"
```


```
g1 <- ggplot(d, aes(x = Comp.1, y = Comp.2,
  label = country)) +
  geom_point() + geom_text_repel() + coord_fixed()
```

- Biplot:

```
g2 <- ggbiplot(track.pc, labels = track$country)
```

## xxx Principal components plot

g1




## xxx Comments on principal components plot

- Good running countries at left of plot: US, UK, Italy, Russia, East and West Germany.
- Bad running countries at right: Western Samoa, Cook Islands.
- Better sprinting countries at bottom: US, Italy, Russia, Brazil, Greece. do is Dominican Republic, where sprinting records relatively good, distance records very bad.
- Better distance-running countries at top: Portugal, Norway, Turkey, Ireland, New Zealand, Mexico. ke is Kenya.

## xxx Biplot

g2



## xxx Comments on biplot

- Had to do some pre-work to interpret PC plot. Biplot more self-contained.
- All variable arrows point right; countries on right have large (bad) record times overall, countries on left good overall.
- Imagine that variable arrows extend negatively as well. Bottom right = bad at distance running, top left = good at distance running.
- Top right = bad at sprinting, bottom left = good at sprinting.
- Doesn't require so much pre-interpretation of components.



# xxx Best 8 running countries

Need to look up two-letter abbreviations in ISO table:

xxx

```
d %>%
  arrange(Comp.1) %>%
  left_join(iso, by = c("country" = "ISO2")) %>%
  select(Comp.1, country, Country) %>%
  slice(1:8)
```

##	Comp.1	country	Country
## 1	-3.462175	us	United States of America
## 2	-3.052104	uk	United Kingdom
## 3	-2.752084	it	Italy
## 4	-2.651062	ru	Russian Federation
## 5	-2.613964	dee	East Germany
## 6	-2.576272	dew	West Germany
## 7	-2.468919	au	Australia
## 8	-2.191917	fr	France

## xxx Worst 8 running countries

```
d %>%
  arrange(desc(Comp.1)) %>%
  left_join(iso, by = c("country" = "ISO2")) %>%
  select(Comp.1, country, Country) %>%
  slice(1:8)
```

##	Comp.1	country	Country
## 1	10.652914	ck	Cook Islands
## 2	7.297865	ws	Samoa
## 3	4.297909	mt	Malta
## 4	3.945224	pg	Papua New Guinea
## 5	3.150886	sg	Singapore
## 6	2.787273	th	Thailand
## 7	2.773125	id	Indonesia
## 8	2.697066	gu	Guam

## xxx Better at distance running

```
d %>%
  arrange(desc(Comp.2)) %>%
  left_join(iso, by = c("country" = "ISO2")) %>%
  select(Comp.2, country, Country) %>%
  slice(1:10)
```

##	Comp.2	country	Country
## 1	1.6860391	cr	Costa Rica
## 2	1.5791490	kp	Korea (North)
## 3	1.5226742	ck	Cook Islands
## 4	1.3957839	tr	Turkey
## 5	1.3167578	pt	Portugal
## 6	1.2829272	gu	Guam
## 7	1.0663756	no	Norway
## 8	0.9547437	ir	Iran, Islamic Republic of
## 9	0.9318729	nz	New Zealand
## 10	0.8495104	mx	Mexico

## xxx Better at sprinting

```
d %>%
  arrange(Comp.2) %>%
  left_join(iso, by = c("country" = "ISO2")) %>%
  select(Comp.2, country, Country) %>%
  slice(1:10)
```

##	Comp.2	country	Country
## 1	-2.4715736	do	Dominican Republic
## 2	-1.9196130	ws	Samoa
## 3	-1.8055052	sg	Singapore
## 4	-1.7832229	bm	Bermuda
## 5	-1.7386063	my	Malaysia
## 6	-1.6851772	th	Thailand
## 7	-1.1204235	us	United States of America
## 8	-0.9989821	it	Italy
## 9	-0.7639385	ru	Russian Federation
## 10	-0.6470634	br	Brazil


# Plot with country names

```
g <- d %>%  
  left_join(iso, by = c("country" = "ISO2")) %>%  
  select(Comp.1, Comp.2, Country) %>%  
  ggplot(aes(x = Comp.1, y = Comp.2, label = Country)) +  
  geom_point() + geom_text_repel(size = 1) +  
  coord_fixed()
```

```
## Warning: Column `country`/`ISO2` joining factor and character  
## vector, coercing into character vector
```

# The plot

g



# xxx Principal components from correlation matrix

Create data file like this:

```
1          0.9705 -0.9600
0.9705     1          -0.9980
-0.9600  -0.9980     1
```

and read in like this:

```
my_url <- "http://www.utsc.utoronto.ca/~butler/d29/cov.txt"
mat <- read_table(my_url, col_names = F)
mat
```

```
## # A tibble: 3 x 3
##       X1      X2      X3
##   <dbl> <dbl> <dbl>
## 1  1      0.970 -0.96
## 2  0.970  1      -0.998
## 3 -0.96  -0.998  1
```

# Pre-processing

A little pre-processing required:


- Turn into matrix (from data frame)
- Feed into princomp as covmat=

```
mat.pc <- mat %>%  
  as.matrix() %>%  
  princomp(covmat = .)
```



# Scree plot: one component fine

```
ggscreeplot(mat.pc)
```



# xxx Component loadings

Compare correlation matrix:

```
mat
```

```
## # A tibble: 3 x 3
##       X1      X2      X3
##   <dbl> <dbl> <dbl>
## 1  1      0.970 -0.96
## 2  0.970  1      -0.998
## 3 -0.96  -0.998  1
```

with component loadings

```
mat.pc$loadings
```

```
##
## Loadings:
##      Comp.1 Comp.2 Comp.3
## X1  0.573  0.812  0.112
## X2  0.581 -0.306 -0.755
## X3 -0.578  0.498 -0.646
##
##
##      Comp.1 Comp.2 Comp.3
## SS loadings  1.000  1.000  1.000
```

# Comments xxx for sign

- When  $X_1$  large,  $X_2$  also large,  $X_3$  small.
- Then comp.1 *negative*.
- When  $X_1$  small,  $X_2$  small,  $X_3$  large.
- Then comp.1 *positive*.

## xxx No scores

- With correlation matrix rather than data, no component scores
- So no principal component plot
- and no biplot.