

STAD29: Statistics for the Life and Social Sciences

Lecture notes

Section 1

Time Series

Packages

Uses my package `mkac` which is on Github. Install with:

```
library(devtools)
install_github("nxskok/mkac")
```

Plus these. You might need to install some of them first: xxx

```
library(ggfortify)
library(forecast)
library(tidyverse)
library(mkac)
```

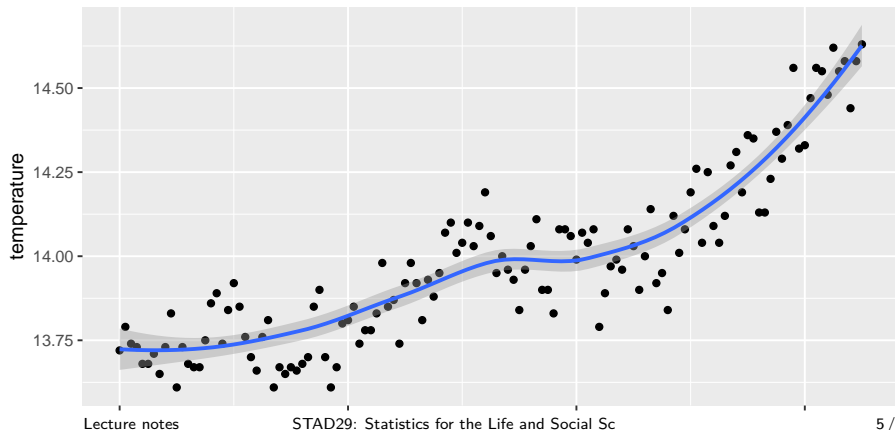
Time trends

- Assess existence or nature of time trends with:
 - correlation
 - regression ideas.
 - (later) time series analysis

World mean temperatures

Global mean temperature every year since 1880: xxx

```
temp=read_csv("temperature.csv")  
ggplot(temp, aes(x=year, y=temperature)) +  
  geom_point() + geom_smooth()
```



Examining trend

- Temperatures increasing on average over time, but pattern very irregular.
- Find (Pearson) correlation with time, and test for significance:

```
with(temp, cor.test(temperature,year))
```

```
##
## Pearson's product-moment correlation
##
## data: temperature and year
## t = 19.996, df = 129, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.8203548 0.9059362
## sample estimates:
## cor
## 0.8695276
```

Comments

- Correlation, 0.8695, significantly different from zero.
- CI shows how far from zero it is.

Tests for *linear* trend with *normal* data.

Kendall correlation

Alternative, Kendall (rank) correlation, which just tests for monotone trend (anything upward, anything downward) and is resistant to outliers:

```
with(temp, cor.test(temperature,year,method="kendall"))
```

```
##
## Kendall's rank correlation tau
##
## data: temperature and year
## z = 11.776, p-value < 2.2e-16
## alternative hypothesis: true tau is not equal to 0
## sample estimates:
##      tau
## 0.6992574
```

Kendall correlation usually closer to 0 for same data, but here P-values comparable. Trend again strongly significant.

Mann-Kendall

- Another way is via **Mann-Kendall**: Kendall correlation with time.
- Use my package `mkac`:

```
kendall_Z_adjusted(temp$temperature)
```

```
## $z
## [1] 11.77267
##
## $z_star
## [1] 4.475666
##
## $ratio
## [1] 6.918858
##
## $P_value
## [1] 0
##
## $P_value_adj
## [1] 7.617357e-06
```

Comments xxx

- Standard Mann-Kendall assumes observations *independent*.
- Observations close together in time often *correlated* with each other.
- Correlation of time series “with itself” called **autocorrelation**.
- Adjusted P-value above is correction for autocorrelation.

Examining rate of change

- Having seen that there *is* a change, question is “how fast is it?”
- Examine slopes:
 - regular regression slope, if you believe straight-line regression
 - Theil-Sen slope: resistant to outliers, based on medians

Ordinary regression against time xxx

```
temp.lm=lm(temperature~year, data=temp)
tidy(temp.lm)
```

```
## # A tibble: 2 x 5
##   term          estimate std.error statistic  p.value
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>
## 1 (Intercept)    2.58      0.570      4.52 1.37e- 5
## 2 year          0.00586   0.000293    20.0 2.42e-41
```

Slope about 0.006 degrees per year (about this many degrees over course of data):

```
coef(temp.lm)[2]*130
```

```
##      year
## 0.7622068
```

Theil-Sen slope

also from mkac:

```
theil_sen_slope(temp$temperature)
```

```
## [1] 0.005675676
```

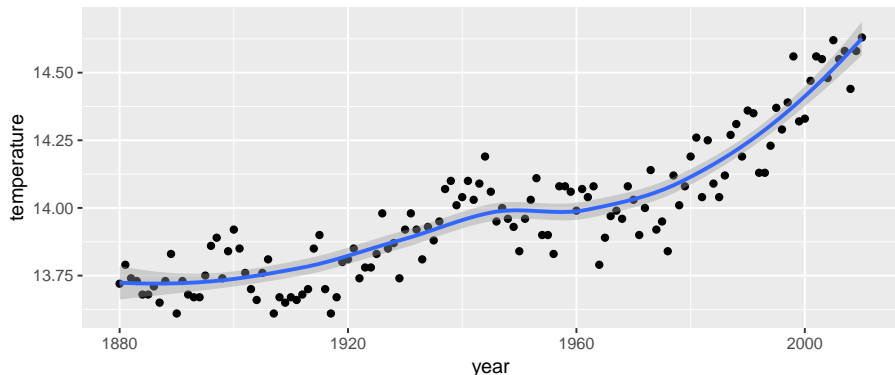
Conclusions

- Slopes:
 - Linear regression: 0.005863
 - Theil-Sen slope: 0.005676
 - Very close.
- Correlations:
 - Pearson 0.8675
 - Kendall 0.6993
 - Kendall correlation smaller, but P-value equally significant (often the case)

Constant rate of change? xxx

Slope assumes that the rate of change is same over all years, but trend seemed to be accelerating: xxx

```
ggplot(temp, aes(x=year, y=temperature)) +  
  geom_point() + geom_smooth()
```



Pre-1970 and post-1970:

```
temp %>%
  mutate(time_period=
    ifelse(year<=1970, "pre-1970", "post-1970")) %>%
  nest(-time_period) %>%
  mutate(theil_sen=map_dbl(
    data, ~theil_sen_slope(.$temperature)))
```

```
## # A tibble: 2 x 3
##   time_period data          theil_sen
##   <chr>         <list>         <dbl>
## 1 pre-1970     <tibble [91 × 4]>    0.00429
## 2 post-1970   <tibble [40 × 4]>    0.0168
```

Theil-Sen slope is very nearly *four times* as big since 1970 vs. before.

Actual time series: the Kings of England

- Age at death of Kings and Queens of England since William the Conqueror (1066):

```
kings=read_table("kings.txt", col_names=F)
```

```
## Parsed with column specification:  
## cols(  
##   X1 = col_double()  
## )
```

Data in one long column X1, so kings is data frame with one column.

Turn into ts time series object

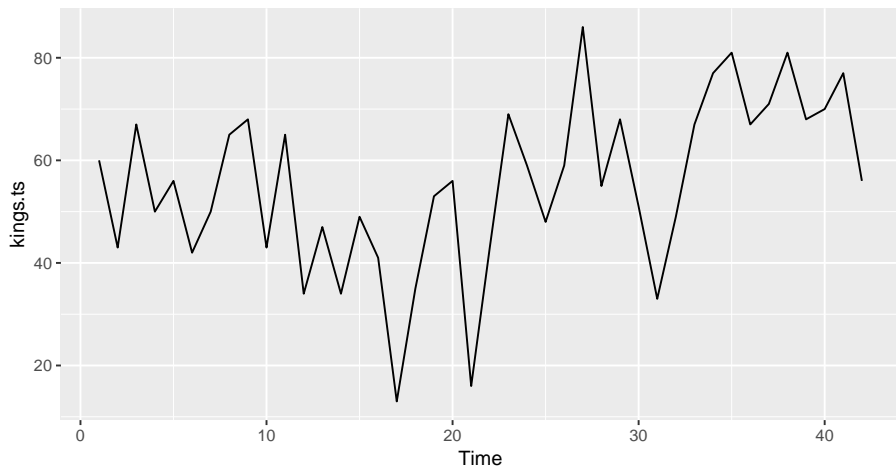
```
kings.ts=ts(kings)
kings.ts
```

```
## Time Series:
## Start = 1
## End = 42
## Frequency = 1
##          X1
## [1,] 60
## [2,] 43
## [3,] 67
## [4,] 50
## [5,] 56
## [6,] 42
## [7,] 50
## [8,] 65
## [9,] 60
```

Plotting a time series xxx

autoplot from ggfortify gives time plot:

```
autoplot(kings.ts)
```



Comments

- “Time” here is order of monarch from William the Conqueror (1st) to George VI (last).
- Looks to be slightly increasing trend of age-at-death
- but lots of irregularity.

Stationarity

A time series is **stationary** if:

- mean is constant over time
- variability constant over time and not changing with mean.

Kings time series seems to have:

- non-constant mean
- but constant variability
- not stationary.

xxx Getting it stationary

- Usual fix for non-stationarity is *differencing*: get new series from original one's values: 2nd - 1st, 3rd - 2nd etc.

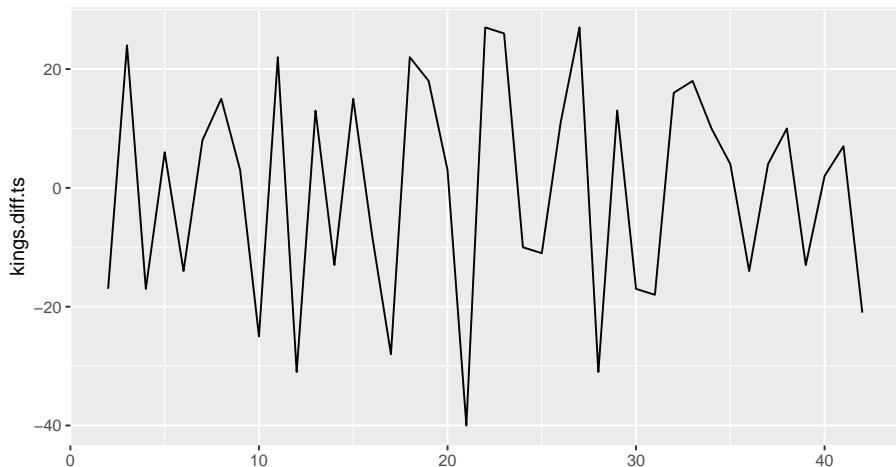
In R, diff:

```
kings.diff.ts=diff(kings.ts)
```

xxx Did differencing fix stationarity?

Looks stationary now: xxx

```
autoplot(kings.diff.ts)
```



xxx Births per month in New York City

from January 1946 to December 1959:

```
ny=read_table("nybirths.txt",col_names=F)
ny
```

```
## # A tibble: 168 x 1
```

```
##       X1
```

```
##    <dbl>
```

```
##  1  26.7
```

```
##  2  23.6
```

```
##  3  26.9
```

```
##  4  24.7
```

```
##  5  25.8
```

```
##  6  24.4
```

```
##  7  24.5
```

```
##  8  23.9
```

```
##  9  23.2
```

```
## 10  23.2
```

```
## # ... with 158 more rows
```


As a time series xxx


```
ny.ts=ts(ny,freq=12,start=c(1946,1))
ny.ts
```

##		Jan	Feb	Mar	Apr	May	Jun
##	1946	26.663	23.598	26.931	24.740	25.806	24.364
##	1947	21.439	21.089	23.709	21.669	21.752	20.761
##	1948	21.937	20.035	23.590	21.672	22.222	22.123
##	1949	21.548	20.000	22.424	20.615	21.761	22.874
##	1950	22.604	20.894	24.677	23.673	25.320	23.583
##	1951	23.287	23.049	25.076	24.037	24.430	24.667
##	1952	23.798	22.270	24.775	22.646	23.988	24.737
##	1953	24.364	22.644	25.565	24.062	25.431	24.635
##	1954	24.657	23.304	26.982	26.199	27.210	26.122
##	1955	24.990	24.239	26.721	23.475	24.767	26.219
##	1956	26.217	24.218	27.914	26.975	28.527	27.139
##	1957	26.589	24.848	27.543	26.896	28.878	27.390
##	1958	27.480	24.884	28.888	26.500	27.884	28.888

xxx Time plot

- Time plot shows extra pattern: xxx

```
autoplot(ny.ts)
```




xxx Comments on time plot

- steady increase (after initial drop)
- repeating pattern each year (seasonal component).
- Not stationary.

xxx Differencing the New York births

Does differencing help here? Looks stationary, but some regular spikes: xxx


```
ny.diff.ts=diff(ny.ts)  
autoplot(ny.diff.ts)
```



xxx Decomposing a seasonal time series

A visual (using original data): xxx

```
ny.d <- decompose(ny.ts)
ny.d %>% autoplot()
```



xxx Decomposition bits

Shows:

- original series
- a “seasonal” part: something that repeats every year
- just the trend, going steadily up (except at the start)
- random: what is left over (“remainder”)

xxx The seasonal part

Fitted seasonal part is same every year, births lowest in February and highest in July:


```
ny.d$seasonal
```

##		Jan	Feb	Mar	Apr
## 1946	-0.6771947	-2.0829607	0.8625232	-0.8016787	
## 1947	-0.6771947	-2.0829607	0.8625232	-0.8016787	
## 1948	-0.6771947	-2.0829607	0.8625232	-0.8016787	
## 1949	-0.6771947	-2.0829607	0.8625232	-0.8016787	
## 1950	-0.6771947	-2.0829607	0.8625232	-0.8016787	
## 1951	-0.6771947	-2.0829607	0.8625232	-0.8016787	
## 1952	-0.6771947	-2.0829607	0.8625232	-0.8016787	
## 1953	-0.6771947	-2.0829607	0.8625232	-0.8016787	
## 1954	-0.6771947	-2.0829607	0.8625232	-0.8016787	
## 1955	-0.6771947	-2.0829607	0.8625232	-0.8016787	
## 1956	-0.6771947	-2.0829607	0.8625232	-0.8016787	

xxx Time series basics: white noise

Each value independent random normal. Knowing one value tells you nothing about the next. “Random” process. xxx

```
wn=rnorm(100)
wn.ts=ts(wn)
autoplot(wn.ts)
```



Lagging a time series

This means moving a time series one (or more) steps back in time:


```
x=rnorm(5)
tibble(x) %>% mutate(x_lagged=lag(x)) -> with_lagged
with_lagged
```

```
## # A tibble: 5 x 2
##       x x_lagged
##   <dbl>   <dbl>
## 1 -2.04    NA
## 2 -0.579 -2.04
## 3  0.608 -0.579
## 4  0.118  0.608
## 5  0.0563 0.118
```

Gain a missing because there is nothing before the first observation.

xxx Lagging white noise

```
tibble(wn) %>% mutate(wn_lagged=lag(wn)) -> wn_with_lagged
ggplot(wn_with_lagged, aes(y=wn, x=wn_lagged))+geom_point()
```



xxx Correlation with lagged series

If you know about white noise at one time point, you know *nothing* about it at the next. This is shown by the scatterplot and the correlation.


On the other hand, this:

```
tibble(age=kings$X1) %>%
  mutate(age_lagged=lag(age)) -> kings_with_lagged
with(kings_with_lagged, cor.test(age, age_lagged))
```

```
##
## Pearson's product-moment correlation
##
## data: age and age_lagged
## t = 2.7336, df = 39, p-value = 0.00937
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.1064770 0.6308209
## sample estimates:
## cor
```

xxx Correlation with next value?

```
ggplot(kings_with_lagged, aes(x=age_lagged, y=age)) +  
  geom_point()
```



xxx Two steps back:

```
kings_with_lagged %>%
  mutate(age_lag_2=lag(age_lagged)) %>%
  with(., cor.test(age, age_lag_2))
```

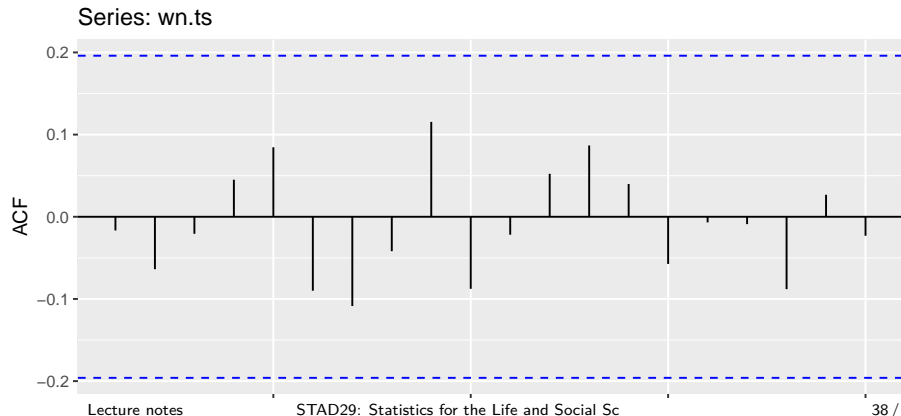
```
##
## Pearson's product-moment correlation
##
## data: age and age_lag_2
## t = 1.5623, df = 38, p-value = 0.1265
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## -0.07128917 0.51757510
## sample estimates:
## cor
## 0.245676
```

Still a correlation two steps back, but smaller (and no longer significant).

xxx Autocorrelation


Correlation of time series with *itself* one, two,... time steps back is useful idea, called **autocorrelation**. Make a plot of it with `acf` and `autoplot`. Here, white noise: xxx

```
acf(wn.ts, plot=F) %>% autoplot()
```



Kings, differenced

```
acf(kings.diff.ts, plot=F) %>% autoplot()
```




xxx Comments on autocorrelations of kings series

Negative autocorrelation at lag 1, nothing beyond that.

- If one value of differenced series positive, next one most likely negative.
- If one monarch lives longer than predecessor, next one likely lives shorter.

NY births, differenced

```
acf(ny.diff.ts, plot=F) %>% autoplot()
```



Lots of stuff:

- large positive autocorrelation at 1.0 years (July one year like July last year)
- large negative autocorrelation at 1 month.
- smallish but significant negative autocorrelation at 0.5 year = 6 months.
- Other stuff – complicated.

xxx Souvenir sales


Monthly sales for a beach souvenir shop in Queensland, Australia:

```
souv=read_table("souvenir.txt", col_names=F)
souv.ts=ts(souv,frequency=12,start=1987)
souv.ts
```

##	Jan	Feb	Mar	Apr
## 1987	1664.81	2397.53	2840.71	3547.29
## 1988	2499.81	5198.24	7225.14	4806.03
## 1989	4717.02	5702.63	9957.58	5304.78
## 1990	5921.10	5814.58	12421.25	6369.77
## 1991	4826.64	6470.23	9638.77	8821.17
## 1992	7615.03	9849.69	14558.40	11587.33
## 1993	10243.24	11266.88	21826.84	17357.33
##	May	Jun	Jul	Aug
## 1987	3752.96	3714.74	4349.61	3566.34
## 1988	5900.88	4951.34	6179.12	4752.15
## 1989	6492.43	6630.80	7349.62	8176.62
## 1990	7609.12	7224.75	8121.22	7979.25
## 1991	8722.37	10209.48	11276.55	12552.22
## 1992	9332.56	13082.09	16732.78	19888.61
## 1993	15997.79	18601.53	26155.15	28586.52
##	Sep	Oct	Nov	Dec
## 1987	5001.80	6400.40	7600.60	10750.01

Plot of souvenir sales

```
autoplot(souv.ts)
```




xxx Several problems:

- Mean goes up over time
- Variability gets larger as mean gets larger
- Not stationary

xxx Problem-fixing:


Fix non-constant variability first by taking logs: xxx

```
souv.log.ts=log(souv.ts)  
autoplot(souv.log.ts)
```



Mean still not constant, so try taking differences

```
souv.log.diff.ts=diff(souv.log.ts)  
autoplot(souv.log.diff.ts)
```




Comments

- Now stationary
- but clear seasonal effect.

Decomposing to see the seasonal effect

```
souv.d=decompose(souv.log.diff.ts)  
autoplot(souv.d)
```



xxx Comments

Big drop in one month's differences. Look at seasonal component to see which:


```
souv.d$seasonal
```

```
##           Jan           Feb           Mar
## 1987      0.23293343  0.49068755
## 1988 -1.90372141  0.23293343  0.49068755
## 1989 -1.90372141  0.23293343  0.49068755
## 1990 -1.90372141  0.23293343  0.49068755
## 1991 -1.90372141  0.23293343  0.49068755
## 1992 -1.90372141  0.23293343  0.49068755
## 1993 -1.90372141  0.23293343  0.49068755
##           Apr           May           Jun
## 1987 -0.39700942  0.02410429  0.05074206
## 1988 -0.39700942  0.02410429  0.05074206
## 1989 -0.39700942  0.02410429  0.05074206
## 1990 -0.39700942  0.02410429  0.05074206
## 1991 -0.39700942  0.02410429  0.05074206
## 1992 -0.39700942  0.02410429  0.05074206
## 1993 -0.39700942  0.02410429  0.05074206
##           Jul           Aug           Sep
## 1987  0.13552988 -0.03710275  0.08650584
## 1988  0.13552988 -0.03710275  0.08650584
## 1989  0.13552988 -0.03710275  0.08650584
## 1990  0.13552988 -0.03710275  0.08650584
## 1991  0.13552988 -0.03710275  0.08650584
## 1992  0.13552988 -0.03710275  0.08650584
## 1993  0.13552988 -0.03710275  0.08650584
##           Oct           Nov           Dec
```

Autocorrelations

```
acf(souv.log.diff.ts, plot=F) %>% autoplot()
```

Series: souv.log.diff.ts



xxx Moving average

- A particular type of time series called a **moving average** or MA process captures idea of autocorrelations at a few lags but not at others.
- Here's generation of MA(1) process, with autocorrelation at lag 1 but not otherwise:

```
beta=1
tibble(e=rnorm(100)) %>%
  mutate(e_lag=lag(e)) %>%
  mutate(y=e+beta*e_lag) %>%
  mutate(y=ifelse(is.na(y), 0, y)) -> ma
```

The series xxx

```
ma
```

```
## # A tibble: 100 x 3
##       e     e_lag     y
##   <dbl>   <dbl> <dbl>
## 1  0.991    NA      0
## 2  0.469    0.991  1.46
## 3  0.535    0.469  1.00
## 4 -0.244    0.535  0.291
## 5  1.17    -0.244  0.928
## 6 -0.473    1.17   0.699
## 7  1.56    -0.473  1.08
## 8 -0.355    1.56   1.20
## 9 -0.400   -0.355 -0.755
## 10 -2.10   -0.400 -2.50
## # ... with 90 more rows
```

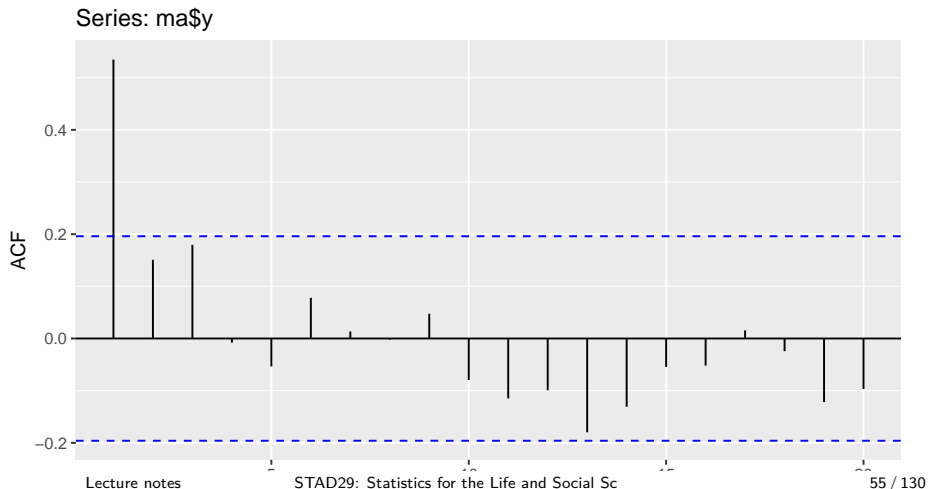
Comments

- e contains independent “random shocks”.
- Start process at 0.
- Then, each value of the time series has that time's random shock, plus a multiple of the last time's random shock.
- $y[i]$ has shock in common with $y[i-1]$; should be a lag 1 autocorrelation.
- But $y[i]$ has no shock in common with $y[i-2]$, so no lag 2 autocorrelation (or beyond).

ACF for MA(1) process xxx

Everything beyond lag 1 appears to be just chance:

```
acf(ma$y, plot=F, na.rm=T) %>% autoplot()
```



xxx AR process

Another kind of time series is AR process, where each value depends on previous one, like this (loop):

```
e=rnorm(100)
x=numeric(0)
x[1]=0
alpha=0.7
for (i in 2:100)
{
  x[i]=alpha*x[i-1]+e[i]
}
```


The series xxx

x


```
##      [1]  0.00000000  0.69150384 -0.27156693
##      [4] -1.69374385 -0.04624706 -0.61289729
##      [7]  0.26464756 -0.21493841 -1.31429232
##     [10]  0.44277420  0.09918044  0.19080999
##     [13] -1.02379326  0.16693770  0.98374525
##     [16]  0.04866219  1.22331904 -0.04784703
##     [19] -0.21367820 -0.68228901  0.25079396
##     [22] -0.86025292  1.75818244  1.19266409
##     [25]  0.30513461  2.41224530  1.28151011
##     [28]  1.68979182  2.01815565  3.53754507
##     [31]  1.85840920  2.32513921  1.77111656
##     [34]  2.12223993  0.91095776  1.58477201
##     [37]  2.08225425  1.09623045 -0.76369221
##     [40] -0.70809836 -1.84439667 -0.38985352
##     [43] -1.04265756 -0.86988314 -1.14485961
```

Comments

- Each random shock now only used for its own value of x
- but $x[i]$ also depends on previous value $x[i-1]$
- so correlated with previous value
- *but* $x[i]$ also contains multiple of $x[i-2]$ and previous x 's
- so all x 's correlated, but autocorrelation dying away.

ACF for AR(1) series

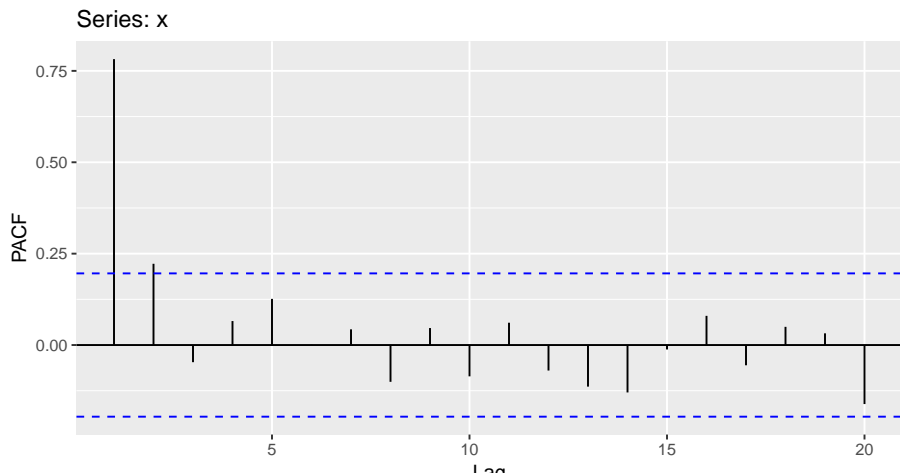
```
acf(x, plot=F) %>% autoplot()
```



xxx Partial autocorrelation function


This cuts off for an AR series: xxx

```
pacf(x, plot=F) %>% autoplot()
```



PACF for an MA series decays slowly

```
pacf(ma$y, plot=F) %>% autoplot()
```



The old way of doing time series analysis

Starting from a series with constant variability (eg. transform first to get it, as for souvenirs):

- Assess stationarity.
- If not stationary, take differences as many times as needed until it is.
- Look at ACF, see if it dies off. If it does, you have MA series.
- Look at PACF, see if that dies off. If it does, have AR series.
- If neither dies off, probably have a mixed “ARMA” series.
- Fit coefficients (like regression slopes).
- Do forecasts.

The new way of doing time series analysis (in R)

- Transform series if needed to get constant variability
- Use package `forecast`.
- Use function `auto.arima` to estimate what kind of series best fits data.
- Use `forecast` to see what will happen in future.

xxx Anatomy of auto.arima output

```
auto.arima(ma$y)
```

```
## Series: ma$y
## ARIMA(0,0,1) with zero mean
##
## Coefficients:
##          ma1
##          0.9070
## s.e.      0.0617
##
## sigma^2 estimated as 0.9878:  log likelihood=-141.64
## AIC=287.29   AICc=287.41   BIC=292.5
```

Comments over.

Comments xxx

- ARIMA part tells you what kind of series you are estimated to have:
 - first number (first 0) is AR (autoregressive) part
 - second number (second 0) is amount of differencing here
 - third number (1) is MA (moving average) part
- Below that, coefficients (with SEs)
- AICc is measure of fit (lower better)

xxx What other models were possible?

Run `auto.arima` with `trace=T`:

```
auto.arima(ma$y, trace=T)
```

```
##
## ARIMA(2,0,2) with non-zero mean : Inf
## ARIMA(0,0,0) with non-zero mean : 345.2328
## ARIMA(1,0,0) with non-zero mean : 313.9535
## ARIMA(0,0,1) with non-zero mean : 287.9463
## ARIMA(0,0,0) with zero mean      : 346.0889
## ARIMA(1,0,1) with non-zero mean : 290.112
## ARIMA(0,0,2) with non-zero mean : 290.1128
## ARIMA(1,0,2) with non-zero mean : 291.7865
## ARIMA(0,0,1) with zero mean      : 287.4124
## ARIMA(1,0,1) with zero mean      : 289.4909
## ARIMA(0,0,2) with zero mean      : 289.4993
## ARIMA(1,0,0) with zero mean      : 312.7625
## ARIMA(1,0,2) with zero mean      : 290.6071
##
```

Doing it all the new way: white noise

```
wn.aa=auto.arima(wn.ts)
wn.aa
```

```
## Series: wn.ts
## ARIMA(0,0,0) with zero mean
##
## sigma^2 estimated as 1.111:  log likelihood=-147.16
## AIC=296.32    AICc=296.36    BIC=298.93
```

Best fit *is* white noise (no AR, no MA, no differencing).

xxx Forecasts:

```
forecast(wn.aa)
```

##	Point Forecast	Lo 80	Hi 80	Lo 95
## 101	0	-1.350869	1.350869	-2.065975
## 102	0	-1.350869	1.350869	-2.065975
## 103	0	-1.350869	1.350869	-2.065975
## 104	0	-1.350869	1.350869	-2.065975
## 105	0	-1.350869	1.350869	-2.065975
## 106	0	-1.350869	1.350869	-2.065975
## 107	0	-1.350869	1.350869	-2.065975
## 108	0	-1.350869	1.350869	-2.065975
## 109	0	-1.350869	1.350869	-2.065975
## 110	0	-1.350869	1.350869	-2.065975

##	Hi 95
## 101	2.065975
## 102	2.065975
## 103	2.065975
## 104	2.065975
## 105	2.065975

MA(1)


```
y.aa=auto.arima(ma$y)
y.aa
```

```
## Series: ma$y
## ARIMA(0,0,1) with zero mean
##
## Coefficients:
##          ma1
##          0.9070
## s.e.    0.0617
##
## sigma^2 estimated as 0.9878:  log likelihood=-141.64
## AIC=287.29   AICc=287.41   BIC=292.5
y.f=forecast(y.aa)
```

Plotting the forecasts for MA(1)

```
autoplot(y.f)
```

Forecasts from ARIMA(0,0,1) with zero mean



xxx AR(1)

```
x.aa=auto.arima(x)
x.aa
```

```
## Series: x
## ARIMA(0,1,1)
##
## Coefficients:
##             ma1
##          -0.3544
## s.e.      0.1062
##
## sigma^2 estimated as 0.979:  log likelihood=-138.99
## AIC=281.97   AICc=282.1   BIC=287.16
```

Oops! Thought it was MA(1), not AR(1)!

xxx Fit right AR(1) model:

```
x.arima=arima(x,order=c(1,0,0))
x.arima
```

```
##
```

```
## Call:
```

```
## arima(x = x, order = c(1, 0, 0))
```

```
##
```

```
## Coefficients:
```

```
##          ar1  intercept
```

```
##          0.7758    -0.3646
```

```
## s.e.    0.0611      0.4220
```


```
##
```

```
## sigma^2 estimated as 0.957:  log likelihood = -140.16,  aic
```


Forecasts for x


```
forecast(x.arima) %>% autoplot()
```

Forecasts from ARIMA(1,0,0) with non-zero mean



Comparing wrong model: xxx

```
forecast(x.aa) %>% autoplot()
```



xxx Kings

```
kings.aa=auto.arima(kings.ts)
kings.aa
```

```
## Series: kings.ts
## ARIMA(0,1,1)
##
## Coefficients:
##             ma1
##          -0.7218
## s.e.      0.1208
##
## sigma^2 estimated as 236.2:  log likelihood=-170.06
## AIC=344.13   AICc=344.44   BIC=347.56
```

xxx Kings forecasts:

```
kings.f=forecast(kings.aa)
kings.f
```

```
##      Point Forecast      Lo 80      Hi 80      Lo 95
## 43      67.75063 48.05479 87.44646 37.62845
## 44      67.75063 47.30662 88.19463 36.48422
## 45      67.75063 46.58489 88.91637 35.38042
## 46      67.75063 45.88696 89.61429 34.31304
## 47      67.75063 45.21064 90.29062 33.27869
## 48      67.75063 44.55402 90.94723 32.27448
## 49      67.75063 43.91549 91.58577 31.29793
## 50      67.75063 43.29362 92.20763 30.34687
## 51      67.75063 42.68718 92.81408 29.41939
## 52      67.75063 42.09507 93.40619 28.51383
##              Hi 95
## 43  97.87281
## 44  99.01703
## 45 100.12084
## 46 101.18822
```

Kings forecasts, plotted

```
autoplot(kings.f) + labs(x="index", y= "age at death")
```

Forecasts from ARIMA(0,1,1)



NY births

Very complicated:

```
ny.aa=auto.arima(ny.ts)
ny.aa
```

```
## Series: ny.ts
## ARIMA(2,1,2)(1,1,1)[12]
##
## Coefficients:
##          ar1          ar2          ma1          ma2          sar1
##          0.6539   -0.4540   -0.7255    0.2532   -0.2427
## s.e.    0.3003    0.2429    0.3227    0.2878    0.0985
##          sma1
##          -0.8451
## s.e.     0.0995
##
## sigma^2 estimated as 0.4076:  log likelihood=-157.45
## AIC=328.91   AICc=329.67   BIC=350.21
```

xxx NY births forecasts


Not *quite* same every year:

```
ny.f=forecast(ny.aa,h=36)
ny.f
```

##		Point Forecast	Lo 80	Hi 80	Lo 95
##	Jan 1960	27.69056	26.87069	28.51043	26.43668
##	Feb 1960	26.07680	24.95838	27.19522	24.36632
##	Mar 1960	29.26544	28.01566	30.51523	27.35406
##	Apr 1960	27.59444	26.26555	28.92333	25.56208
##	May 1960	28.93193	27.52089	30.34298	26.77392
##	Jun 1960	28.55379	27.04381	30.06376	26.24448
##	Jul 1960	29.84713	28.23370	31.46056	27.37960
##	Aug 1960	29.45347	27.74562	31.16132	26.84155
##	Sep 1960	29.16388	27.37259	30.95517	26.42433
##	Oct 1960	29.21343	27.34498	31.08188	26.35588
##	Nov 1960	27.26221	25.31879	29.20563	24.29000
##	Dec 1960	28.06863	26.05137	30.08589	24.98349
##	Jan 1961	27.66908	25.59684	29.74132	24.49986

Plotting the forecasts

```
autoplot(ny.f)+labs(x="time", y="births")
```



Log-souvenir sales

```
souv.aa=auto.arima(souv.log.ts)
souv.aa
```

```
## Series: souv.log.ts
## ARIMA(2,0,0)(0,1,1)[12] with drift
##
## Coefficients:
##          ar1      ar2      sma1    drift
##      0.3470  0.3516  -0.5205  0.0238
## s.e.  0.1092  0.1115   0.1700  0.0031
##
## sigma^2 estimated as 0.02953:  log likelihood=24.54
## AIC=-39.09   AICc=-38.18   BIC=-27.71
```

```
souv.f=forecast(souv.aa,h=27)
```

xxx The forecasts

Differenced series showed low value for January (large drop). December highest, Jan and Feb lowest:


```
souv.f
```

##	Point Forecast	Lo 80	Hi 80
## Jan 1994	9.578291	9.358036	9.798545
## Feb 1994	9.754836	9.521700	9.987972
## Mar 1994	10.286195	10.030937	10.541453
## Apr 1994	10.028630	9.765727	10.291532
## May 1994	9.950862	9.681555	10.220168
## Jun 1994	10.116930	9.844308	10.389551
## Jul 1994	10.369140	10.094251	10.644028
## Aug 1994	10.460050	10.183827	10.736274
## Sep 1994	10.535595	10.258513	10.812677
## Oct 1994	10.585995	10.308386	10.863604
## Nov 1994	11.017734	10.739793	11.295674
## Dec 1994	11.795964	11.517817	12.074111
## Jan 1995	9.840884	9.540241	10.141527
## Feb 1995	10.015540	9.711785	10.319295
## Mar 1995	10.555070	10.246346	10.863794
## Apr 1995	10.299676	9.989043	10.610309
## May 1995	10.225535	9.913326	10.537743

Plotting the forecasts

```
autoplot(souv.f)
```

Forecasts from ARIMA(2,0,0)(0,1,1)[12] with drift



Global mean temperatures, revisited


```
temp.ts=ts(temp$temperature,start=1880)
temp.aa=auto.arima(temp.ts)
temp.aa
```

```
## Series: temp.ts
## ARIMA(1,1,3) with drift
##
## Coefficients:
##          ar1      ma1      ma2      ma3      drift
##      -0.9374  0.5038  -0.6320  -0.2988  0.0067
## s.e.   0.0835  0.1088  0.0876  0.0844  0.0025
##
## sigma^2 estimated as 0.008939:  log likelihood=124.34
## AIC=-236.67   AICc=-235.99   BIC=-219.47
```

Forecasts

```
temp.f=forecast(temp.aa)  
autoplot(temp.f)+labs(x="year", y="temperature")
```

Forecasts from ARIMA(1,1,3) with drift



Section 2

Multiway frequency tables

Packages

```
library(tidyverse)
```

Multi-way frequency analysis

- A study of gender and eyewear-wearing finds the following frequencies:

Gender	Contacts	Glasses	None
Female	121	32	129
Male	42	37	85

- Is there association between eyewear and gender?
- Normally answer this with chisquare test (based on observed and expected frequencies from null hypothesis of no association).
- Two categorical variables and a frequency.
- We assess in way that generalizes to more categorical variables.

The data file

```
gender contacts glasses none
female 121      32      129
male   42      37      85
```

- This is *not tidy*!
- Two variables are gender and *eyewear*, and those numbers all frequencies.

```
my_url <- "http://www.utsc.utoronto.ca/~butler/d29/eyewear.txt"
eyewear <- read_delim(my_url, " ")
eyewear
```

```
## # A tibble: 2 x 4
##   gender contacts glasses none
##   <chr>      <dbl>    <dbl> <dbl>
## 1 female      121      32    129
```

Tidying the data

```
eyes <- eyewear %>%
  gather(eyewear, frequency, contacts:none)
eyes
```

```
## # A tibble: 6 x 3
##   gender eyewear frequency
##   <chr>   <chr>      <dbl>
## 1 female contacts      121
## 2 male   contacts      42
## 3 female glasses      32
## 4 male   glasses      37
## 5 female none        129
## 6 male   none         85
```

```
xt <- xtabs(frequency ~ gender + eyewear, data = eyes)
xt
```

Modelling

- Last table on previous page is “reconstituted” contingency table, for checking.
- Predict frequency from other factors and combos. glm with poisson family.

```
eyes.1 <- glm(frequency ~ gender * eyewear,  
  data = eyes,  
  family = "poisson"  
)
```

def

- Called **log-linear model**.

What can we get rid of?

```
{
drop1(eyes.1, test = "Chisq")

## Single term deletions
##
## Model:
## frequency ~ gender * eyewear
##              Df Deviance      AIC      LRT  Pr(>Chi)
## <none>                0.000 47.958
## gender:eyewear  2    17.829 61.787 17.829 0.0001345
##
## <none>
## gender:eyewear ***
## ---
## Signif. codes:
## 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

def }
```

Conclusions

- drop1 says what we can remove at this step. Significant = must stay.
- Cannot remove anything.
- Frequency depends on gender-wear *combination*, cannot be simplified further.
- Gender and eyewear are *associated*.
- Stop here.

prop.table

Original table:

```
{
xt

##           eyewear
## gender  contacts glasses none
## female    121      32  129
## male      42      37   85
```

} Calculate eg. row proportions like this:

```
{
prop.table(xt, margin = 1)

##           eyewear
## gender  contacts glasses      none
## female 0.4290780 0.1134752 0.4574468
## male  0.2560976 0.2256098 0.5182927
```

No association

- Suppose table had been as shown below:

```
my_url <- "http://www.utsc.utoronto.ca/~butler/d29/eyewear2.txt"
eyewear2 <- read_table(my_url)
eyes2 <- eyewear2 %>% gather(eyewear, frequency, contacts:none)
xt2 <- xtabs(frequency ~ gender + eyewear, data = eyes2)
xt2
```

```
##           eyewear
## gender  contacts glasses none
##  female      150      30  120
##   male       75      16   62
```

```
prop.table(xt2, margin = 1)
```

```
##           eyewear
## gender  contacts glasses      none
##  female 0.5000000 0.1000000 0.4000000
```

Analysis for revised data

```
eyes.2 <- glm(frequency ~ gender * eyewear,
  data = eyes2,
  family = "poisson"
)
drop1(eyes.2, test = "Chisq")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## frequency ~ gender * eyewear
```

##		Df	Deviance	AIC	LRT	Pr(>Chi)
##	<none>		0.000000	47.467		
##	gender:eyewear	2	0.047323	43.515	0.047323	0.9766

No longer any association. Take out interaction.

No interaction

```
{
eyes.3 <- update(eyes.2, . ~ . - gender:eyewear)
drop1(eyes.3, test = "Chisq")

## Single term deletions
##
## Model:
## frequency ~ gender + eyewear
##           Df Deviance      AIC      LRT  Pr(>Chi)
## <none>           0.047   43.515
## gender      1   48.624   90.091   48.577 3.176e-12 ***
## eyewear     2  138.130  177.598  138.083 < 2.2e-16 ***
## ---
## Signif. codes:
## 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
}
```

Chest pain, being overweight and being a smoker

- In a hospital emergency department, 176 subjects who attended for acute chest pain took part in a study.
- Each subject had a normal or abnormal electrocardiogram reading (ECG), were overweight (as judged by BMI) or not, and were a smoker or not.
- How are these three variables related, or not?

The data

In modelling-friendly format:

```
ecg bmi smoke count  
abnormal overweight yes 47  
abnormal overweight no 10  
abnormal normalweight yes 8  
abnormal normalweight no 6  
normal overweight yes 25  
normal overweight no 15  
normal normalweight yes 35  
normal normalweight no 30
```

First step

```
my_url <- "http://www.utsc.utoronto.ca/~butler/d29/ecg.txt"
chest <- read_delim(my_url, " ")
chest.1 <- glm(count ~ ecg * bmi * smoke,
  data = chest,
  family = "poisson"
)
drop1(chest.1, test = "Chisq")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## count ~ ecg * bmi * smoke
```

```
##           Df Deviance      AIC      LRT Pr(>Chi)
```

```
## <none>           0.0000 53.707
```

```
## ecg:bmi:smoke   1    1.3885 53.096 1.3885    0.2387
```

That 3-way interaction comes out.

Removing the 3-way interaction

```
chest.2 <- update(chest.1, . ~ . - ecg:bmi:smoke)
drop1(chest.2, test = "Chisq")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## count ~ ecg + bmi + smoke + ecg:bmi + ecg:smoke + bmi:smoke
```

```
##           Df Deviance      AIC      LRT  Pr(>Chi)
```

```
## <none>           1.3885 53.096
```

```
## ecg:bmi      1  29.0195 78.727 27.6310 1.468e-07 ***
```

```
## ecg:smoke    1   4.8935 54.601  3.5050  0.06119 .
```

```
## bmi:smoke    1   4.4689 54.176  3.0803  0.07924 .
```

```
## ---
```

```
## Signif. codes:
```

```
## 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

At $\alpha = 0.05$, bmi:smoke comes out.

Removing bmi:smoke

```
chest.3 <- update(chest.2, . ~ . - bmi:smoke)
drop1(chest.3, test = "Chisq")
```

```
## Single term deletions
##
## Model:
## count ~ ecg + bmi + smoke + ecg:bmi + ecg:smoke
##           Df Deviance    AIC    LRT  Pr(>Chi)
## <none>           4.469 54.176
## ecg:bmi      1   36.562 84.270 32.094 1.469e-08 ***
## ecg:smoke    1   12.436 60.144  7.968 0.004762 **
## ---
## Signif. codes:
## 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

ecg:smoke has become significant. So we have to stop.

Understanding the final model

- Thinking of `ecg` as “response” that might depend on anything else.
- What is associated with `ecg`? Both `bmi` on its own and `smoke` on its own, but *not* the combination of both.
- `ecg:bmi` table:

```
xtabs(count ~ ecg + bmi, data = chest)
```

##		bmi	
##	ecg	normalweight	overweight
##	abnormal	14	57
##	normal	65	40

- Most normal weight people have a normal ECG, but a majority of overweight people have an *abnormal* ECG. That is, knowing about BMI says something about likely ECG.

ecg:smoke

- ecg:smoke table:

```
xtabs(count ~ ecg + smoke, data = chest)
```

```
##           smoke
## ecg         no  yes
##  abnormal  16   55
##   normal   45   60
```

- Most nonsmokers have a normal ECG, but smokers are about 50–50 normal and abnormal ECG.
- Don't look at smoke:bmi table since not significant.

Simpson's paradox: the airlines example

Airport	Alaska Airlines		America West	
	On time	Delayed	On time	Delayed
Los Angeles	497	62	694	117
Phoenix	221	12	4840	415
San Diego	212	20	383	65
San Francisco	503	102	320	129
Seattle	1841	305	201	61
Total	3274	501	6438	787

Use status as variable name for “on time/delayed”.

- Alaska: 13.3% flights delayed ($501/(3274 + 501)$).
- America West: 10.9% ($787/(6438 + 787)$).
- America West more punctual, right?

Arranging the data

- Can only have single thing in columns, so we have to construct column names like this: `\begin{small}`

airport	aa_ontime	aa_delayed	aw_ontime	aw_delayed
LosAngeles	497	62	694	117
Phoenix	221	12	4840	415
SanDiego	212	20	383	65
SanFrancisco	503	102	320	129
Seattle	1841	305	201	61

`\end{small}`

- Some tidying gets us the right layout, with frequencies all in one column and the airline and delayed/on time status separated out:

```
my_url <- "http://www.utsc.utoronto.ca/~butler/d29/airlines.txt"
airlines <- read_table2(my_url)
```

The data frame punctual

```
## # A tibble: 20 x 4
##   airport      airline status   freq
##   <chr>        <chr>   <chr> <dbl>
## 1 LosAngeles   aa      ontime  497
## 2 Phoenix      aa      ontime  221
## 3 SanDiego     aa      ontime  212
## 4 SanFrancisco aa      ontime  503
## 5 Seattle      aa      ontime 1841
## 6 LosAngeles   aa      delayed  62
## 7 Phoenix      aa      delayed  12
## 8 SanDiego     aa      delayed  20
## 9 SanFrancisco aa      delayed 102
## 10 Seattle     aa      delayed 305
## 11 LosAngeles  aw      ontime  694
## 12 Phoenix     aw      ontime 4840
## 13 SanDiego    aw      ontime  383
```

Proportions delayed by airline

- Two-step process: get appropriate subtable:

```
xt <- xtabs(freq ~ airline + status, data = punctual)
xt
```

```
##           status
## airline delayed ontime
##      aa      501   3274
##      aw      787   6438
```

- and then calculate appropriate proportions:

```
prop.table(xt, margin = 1)
```

```
##           status
## airline  delayed   ontime
##      aa 0.1327152 0.8672848
##      aw 0.1089273 0.8910727
```

Proportion delayed by airport, for each airline

```
xt <- xtabs(freq ~ airline + status + airport, data = punctual)
xp <- prop.table(xt, margin = c(1, 3))
ftable(xp,
  row.vars = c("airport", "airline"),
  col.vars = "status"
)
```

##		status	delayed	ontime
##	airport	airline		
##	LosAngeles	aa	0.11091234	0.88908766
##		aw	0.14426634	0.85573366
##	Phoenix	aa	0.05150215	0.94849785
##		aw	0.07897241	0.92102759
##	SanDiego	aa	0.08620690	0.91379310
##		aw	0.14508929	0.85491071
##	SanFrancisco	aa	0.16859504	0.83140496
##		aw	0.08720512	0.91279488

Simpson's Paradox

Airport	Alaska	America West
Los Angeles	11.4	14.4
Phoenix	5.2	7.9
San Diego	8.6	14.5
San Francisco	16.9	28.7
Seattle	14.2	23.2
Total	13.3	10.9

- America West more punctual overall,
- but worse at *every single* airport!
- How is that possible?
- Log-linear analysis sheds some light.

Model 1 and output

```
punctual.1 <- glm(freq ~ airport * airline * status,
  data = punctual, family = "poisson"
)
drop1(punctual.1, test = "Chisq")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## freq ~ airport * airline * status
```

```
##                Df Deviance    AIC    LRT
```

```
## <none>                0.0000 183.44
```

```
## airport:airline:status  4    3.2166 178.65 3.2166
```

```
##                Pr(>Chi)
```

```
## <none>
```

```
## airport:airline:status  0.5223
```

```
def
```

Remove 3-way interaction

```
punctual.2 <- update(punctual.1, ~ . - airport:airline:status)
drop1(punctual.2, test = "Chisq")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## freq ~ airport + airline + status + airport:airline + airport:status
```

```
##      airline:status
```

```
##              Df Deviance      AIC      LRT  Pr(>Chi)
```

```
## <none>              3.2   178.7
```

```
## airport:airline    4   6432.5 6599.9 6429.2 < 2.2e-16
```

```
## airport:status     4    240.1  407.5  236.9 < 2.2e-16
```

```
## airline:status     1     45.5  218.9   42.2 8.038e-11
```

```
##
```

```
## <none>
```

```
## airport:airline ***
```

```
""
```


Understanding the significance

- `airline:status:`

```
xt <- xtabs(freq ~ airline + status, data = punctual)
prop.table(xt, margin = 1)
```

```
##           status
## airline  delayed   ontime
##      aa 0.1327152 0.8672848
##      aw 0.1089273 0.8910727
```

- More of Alaska Airlines' flights delayed overall.
- Saw this before.

Understanding the significance (2)

- airport:status:

```
xt <- xtabs(freq ~ airport + status, data = punctual)
prop.table(xt, margin = 1)
```

##		status	
##	airport	delayed	ontime
##	LosAngeles	0.13065693	0.86934307
##	Phoenix	0.07780612	0.92219388
##	SanDiego	0.12500000	0.87500000
##	SanFrancisco	0.21916509	0.78083491
##	Seattle	0.15199336	0.84800664

- Flights into San Francisco (and maybe Seattle) are often late, and flights into Phoenix are usually on time.
- Considerable variation among airports.

Understanding the significance (3)

- airport:airline:

```
xt <- xtabs(freq ~ airport + airline, data = punctual)
prop.table(xt, margin = 2)
```

```
##                airline
## airport          aa          aw
##  LosAngeles    0.14807947 0.11224913
##   Phoenix      0.06172185 0.72733564
##  SanDiego       0.06145695 0.06200692
##  SanFrancisco  0.16026490 0.06214533
##   Seattle      0.56847682 0.03626298
```

- What fraction of each airline's flights are to each airport.
- Most of Alaska Airlines' flights to Seattle and San Francisco.
- Most of America West's flights to Phoenix.

The resolution

- Most of America West's flights to Phoenix, where it is easy to be on time.
- Most of Alaska Airlines' flights to San Francisco and Seattle, where it is difficult to be on time.
- Overall comparison looks bad for Alaska because of this.
- But, *comparing like with like*, if you compare each airline's performance *to the same airport*, Alaska does better.
- Aggregating over the very different airports was a (big) mistake: that was the cause of the Simpson's paradox.
- Alaska Airlines is *more* punctual when you do the proper comparison.

Ovarian cancer: a four-way table

- Retrospective study of ovarian cancer done in 1973.
- Information about 299 women operated on for ovarian cancer 10 years previously.
- Recorded:
 - stage of cancer (early or advanced)
 - type of operation (radical or limited)
 - X-ray treatment received (yes or no)
 - 10-year survival (yes or no)
 - Survival looks like response (suggests logistic regression).
 - Log-linear model finds any associations at all.

The data

after tidying:

{

```
stage operation xray survival freq
early radical no no 10
early radical no yes 41
early radical yes no 17
early radical yes yes 64
early limited no no 1
early limited no yes 13
early limited yes no 3
early limited yes yes 9
advanced radical no no 38
advanced radical no yes 6
advanced radical yes no 64
advanced radical yes yes 11
advanced limited no no 3
advanced limited no yes 1
advanced limited yes no 13
advanced limited yes yes 5
```

Stage 1

hopefully looking familiar by now:

```
my_url <- "http://www.utsc.utoronto.ca/~butler/d29/cancer.txt"
cancer <- read_delim(my_url, " ")
cancer %>% print(n = 6)
```

```
## # A tibble: 16 x 5
##   stage operation xray  survival  freq
##   <chr> <chr>      <chr> <chr>      <dbl>
## 1 early radical   no    no         10
## 2 early radical   no    yes        41
## 3 early radical   yes   no         17
## 4 early radical   yes   yes        64
## 5 early limited   no    no          1
## 6 early limited   no    yes        13
## # ... with 10 more rows
```

```
cancer_1 <- glm(freq ~ stage * operation * xray * survival
```

Output 1

See what we can remove:

```
drop1(cancer.1, test = "Chisq")
```

```
## Single term deletions
##
## Model:
## freq ~ stage * operation * xray * survival
##
```

	Df	Deviance	AIC
## <none>		0.00000	98.130
## stage:operation:xray:survival	1	0.60266	96.732

```
##
##
```

	LRT	Pr(>Chi)
## <none>		
## stage:operation:xray:survival	0.60266	0.4376

```
def
```

Non-significant interaction can come out.

Stage 2

```
cancer.2 <- update(cancer.1, ~ .
- stage:operation:xray:survival)
drop1(cancer.2, test = "Chisq")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## freq ~ stage + operation + xray + survival + stage:operation
```

```
##      stage:xray + operation:xray + stage:survival + operation
```

```
##      xray:survival + stage:operation:xray + stage:operation
```

```
##      stage:xray:survival + operation:xray:survival
```

```
##                                     Df Deviance    AIC      LRT
```

```
## <none>                                0.60266 96.732
```

```
## stage:operation:xray                   1   2.35759 96.487 1.75493
```

```
## stage:operation:survival               1   1.17730 95.307 0.57465
```

```
## stage:xray:survival                    1   0.95577 95.085 0.35311
```

```
""
```

Take out stage:xray:survival

```
cancer.3 <- update(cancer.2, . ~ . - stage:xray:survival)
drop1(cancer.3, test = "Chisq")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## freq ~ stage + operation + xray + survival + stage:operation
```

```
##      stage:xray + operation:xray + stage:survival + operation:survival
```

```
##      xray:survival + stage:operation:xray + stage:operation:survival
```

```
##      operation:xray:survival
```

```
##              Df Deviance      AIC      LRT
```

```
## <none>              0.95577 95.085
```

```
## stage:operation:xray      1  3.08666 95.216 2.13089
```

```
## stage:operation:survival  1  1.56605 93.696 0.61029
```

```
## operation:xray:survival   1  1.55124 93.681 0.59547
```

```
##              Pr(>Chi)
```

```
""
```

Remove operation:xray:survival

```
cancer.4 <- update(cancer.3, . ~ . - operation:xray:survival)
drop1(cancer.4, test = "Chisq")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## freq ~ stage + operation + xray + survival + stage:operation
```

```
##      stage:xray + operation:xray + stage:survival + operation
```

```
##      xray:survival + stage:operation:xray + stage:operation
```

```
##              Df Deviance    AIC    LRT
```

```
## <none>                1.5512 93.681
```

```
## xray:survival          1   1.6977 91.827 0.1464
```

```
## stage:operation:xray   1   6.8420 96.972 5.2907
```

```
## stage:operation:survival 1   1.9311 92.061 0.3799
```

```
##              Pr(>Chi)
```

```
## <none>
```

```
""
```

Comments

- `stage:operation:xray` has now become significant, so won't remove that.
- Shows value of removing terms one at a time.
- There are no higher-order interactions containing both `xray` and `survival`, so now we get to test (and remove) `xray:survival`.

Remove xray:survival

```
cancer.5 <- update(cancer.4, . ~ . - xray:survival)
drop1(cancer.5, test = "Chisq")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## freq ~ stage + operation + xray + survival + stage:operation
```

```
##      stage:xray + operation:xray + stage:survival + operation:survival
```

```
##      stage:operation:xray + stage:operation:survival
```

```
##                                     Df Deviance    AIC    LRT
```

```
## <none>                                1.6977 91.827
```

```
## stage:operation:xray                1    6.9277 95.057 5.2300
```

```
## stage:operation:survival            1    2.0242 90.154 0.3265
```

```
##                                     Pr(>Chi)
```

```
## <none>
```

```
## stage:operation:xray                0.0222 *
```

```
## stage:operation:survival            0.5677
```

Remove stage:operation:survival

```
cancer.6 <- update(cancer.5, . ~ . - stage:operation:survival)
drop1(cancer.6, test = "Chisq")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## freq ~ stage + operation + xray + survival + stage:operation
```

```
##      stage:xray + operation:xray + stage:survival + operation:survival
```

```
##      stage:operation:xray
```

```
##              Df Deviance      AIC      LRT
```

```
## <none>                2.024   90.154
```

```
## stage:survival         1  135.198  221.327 133.173
```

```
## operation:survival     1    4.116   90.245  2.092
```

```
## stage:operation:xray   1    7.254   93.384  5.230
```

```
##              Pr(>Chi)
```

```
## <none>
```

```
""
```

Last step?

Remove operation:survival.

```
cancer.7 <- update(cancer.6, . ~ . - operation:survival)
drop1(cancer.7, test = "Chisq")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## freq ~ stage + operation + xray + survival + stage:operation
```

```
##      stage:xray + operation:xray + stage:survival + stage:operation
```

```
##              Df Deviance      AIC      LRT
```

```
## <none>                4.116   90.245
```

```
## stage:survival         1  136.729  220.859 132.61
```

```
## stage:operation:xray   1    9.346   93.475   5.23
```

```
##              Pr(>Chi)
```

```
## <none>
```

```
## stage:survival         <2e-16 ***
```

Conclusions

- What matters is things associated with survival (survival is “response”).
- Only significant such term is stage:survival:

```
xt <- xtabs(freq ~ stage + survival, data = cancer)
prop.table(xt, margin = 1)
```

```
##           survival
## stage           no           yes
##   advanced 0.8368794 0.1631206
##   early    0.1962025 0.8037975
```

- Most people in early stage of cancer survived, and most people in advanced stage did not survive.
- This true *regardless* of type of operation or whether or not X-ray treatment was received. These things have no impact on survival.

What about that other interaction?

```
xt <- xtabs(freq ~ operation + xray + stage, data = cancer)
ftable(prop.table(xt, margin = 3))
```

```
##                stage  advanced      early
## operation xray
## limited   no      0.02836879 0.08860759
##           yes      0.12765957 0.07594937
## radical   no      0.31205674 0.32278481
##           yes      0.53191489 0.51265823
```

- Out of the people at each stage of cancer (since margin=3 and stage was listed 3rd).
- The association is between stage and xray *only for those who had the limited operation*.
- For those who had the radical operation, there was no association between stage and xray.

General procedure

- Start with “complete model” including all possible interactions.
- `drop1` gives highest-order interaction(s) remaining, remove least non-significant.
- Repeat as necessary until everything significant.
- Look at subtables of significant interactions.
- Main effects not usually very interesting.
- Interactions with “response” usually of most interest: show association with response.

```
## Error in FUN(X[[i]], ...): invalid 'name' argument
```