


Let me create some random data:

```
> set.seed(457299)
> x=rnorm(100,10,3)
> boxplot(x)
```




This is normal data, so we're not expecting any serious outliers. The boxplot reveals two small outliers, one at each end. A summary of `x` looks like this:

```
> summary(x)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.536	8.026	10.030	9.930	11.780	18.010


These data were generated from a normal distribution, but how normal do they look? One answer is to get a histogram:

```
> hist(x)
```



and another is a normal quantile plot:

```
> qqnorm(x)
> qqline(x)
```



These both look pretty normal, so the suggestion is that those outliers on the boxplot are not to be taken seriously.