

Time Series

Correlation and slopes with time

Time trends

- Assess existence or nature of time trends with:
- correlation
- regression ideas.

World mean temperatures

Global mean temperature every year since 1880:

```
temp = read.csv("temperature.csv")
attach(temp)
plot(temperature ~ year, type = "b")
lines(lowess(temperature ~ year))
```

Examining trend

- Temperatures increasing on average over time, but pattern very irregular.
- Find (Pearson) correlation with time, and test for significance:

```
cor.test(temperature, year)

##
##  Pearson's product-moment correlation
##
## data:  temperature and year
## t = 20, df = 129, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
##  0.8204 0.9059
## sample estimates:
##      cor
## 0.8695
```

- Correlation, 0.8695, significantly different from zero.
- CI shows how far from zero it is.

Tests for *linear* trend with *normal* data.




Figure 1:

Kendall correlation

Alternative, Kendall (rank) correlation, which just tests for monotone trend (anything upward, anything downward) and is resistant to outliers:

```
cor.test(temperature, year, method = "kendall")

##
## Kendall's rank correlation tau
##
## data: temperature and year
## z = 11.78, p-value < 2.2e-16
## alternative hypothesis: true tau is not equal to 0
## sample estimates:
## tau
## 0.6993
```

Kendall correlation usually closer to 0 for same data, but here P-values comparable. Trend again strongly significant.

Mann-Kendall

Another way is via **Mann-Kendall**: Kendall correlation with time:

```
library(Kendall)
MannKendall(temperature)

## tau = 0.699, 2-sided pvalue =<2e-16
```

Answer same as previous.

Examining rate of change

- Having seen that there *is* a change, question is “how fast is it?”
- Examine slopes:
- regular regression slope, if you believe straight-line regression
- Theil-Sen slope: resistant to outliers, based on medians

Ordinary regression against time

```
temp.lm = lm(temperature ~ year)
summary(temp.lm)

##
## Call:
## lm(formula = temperature ~ year)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.3250 -0.1012  0.0058  0.0836  0.2850
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.579420   0.570398   4.52  1.4e-05 ***
## year         0.005863   0.000293  20.00 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.127 on 129 degrees of freedom
## Multiple R-squared:  0.756,    Adjusted R-squared:  0.754
## F-statistic:  400 on 1 and 129 DF,  p-value: <2e-16
```

Slope about 0.006 degrees per year (about 0.8 degrees over course of data).

Theil-Sen slope

Two ways, first more or less like regression:

```
library(zyp)
temp.zs = zyp.sen(temperature ~ year, data = temp)
temp.zs

##
## Call:
## NULL
##
## Coefficients:
## Intercept      year
##  2.94703      0.00568
```

Theil-Sen slope, the second way

Don't need to supply `year`, but lots of output:

```
zyp.trend.vector(temperature)
```

```
##      lbound      trend      trendp      ubound      tau      sig
##  0.005000  0.005676  0.743514  0.006333  0.472925  0.000000
##      nruns      autocor valid_frac      linear intercept
##  1.000000  0.614048  1.000000  0.005863  13.611622
```

```
detach(temp)
```

- `trend`: Theil-Sen slope per unit time (year, here)
- `lbound`, `ubound`: confidence interval (95% by default) for slope
- `trendp`: slope over entire time period
- `linear`: regression slope on same data

Other things confuse me!

Conclusions

- Linear regression slope is 0.005863
- Theil-Sen slope is 0.005676
- Very close.
- Pearson correlation is 0.8675
- Kendall correlation is 0.6993
- Kendall correlation smaller, but P-value equally significant (usually the case)

Actual time series

The Kings of England

- Age at death of Kings and Queens of England since William the Conqueror (1066):

```
kings = read.table("kings.txt", header = F)
kings.ts = ts(kings)
```

Data in one long column V1, so `kings` is data frame with one column. Turn into `ts` time series object.

```
kings.ts
```

```
## Time Series:
## Start = 1
## End = 42
## Frequency = 1
##      V1
## [1,] 60
## [2,] 43
## [3,] 67
## [4,] 50
## [5,] 56
## [6,] 42
## [7,] 50
## [8,] 65
## [9,] 68
## [10,] 43
## [11,] 65
## [12,] 34
## [13,] 47
## [14,] 34
## [15,] 49
## [16,] 41
## [17,] 13
## [18,] 35
## [19,] 53
## [20,] 56
## [21,] 16
## [22,] 43
## [23,] 69
## [24,] 59
## [25,] 48
## [26,] 59
## [27,] 86
## [28,] 55
## [29,] 68
## [30,] 51
## [31,] 33
## [32,] 49
```

```
## [33,] 67
## [34,] 77
## [35,] 81
## [36,] 67
## [37,] 71
## [38,] 81
## [39,] 68
## [40,] 70
## [41,] 77
## [42,] 56
```

Plotting a time series

Plotting gives time plot:

```
plot(kings.ts)
lines(lowess(kings.ts))
```

Comments

- “Time” here is order of monarch from William the Conqueror (1st) to George VI (last).
- Looks to be slightly increasing trend of age-at-death
- but lots of irregularity.

Stationarity

A time series is **stationary** if:

- mean is constant over time
- variability constant over time and not changing with mean.

Kings time series seems to have:

- non-constant mean
- but constant variability
- not stationary.




Figure 2:

Usual fix for non-stationarity is *differencing*: new series 2nd - 1st, 3rd - 2nd etc.

In R, diff:

```
kings.diff.ts = diff(kings.ts)
```

Did differencing fix stationarity?

Looks stationary now:

```
plot(kings.diff.ts, main = "Kings series, differenced")  
lines(lowess(kings.diff.ts))
```

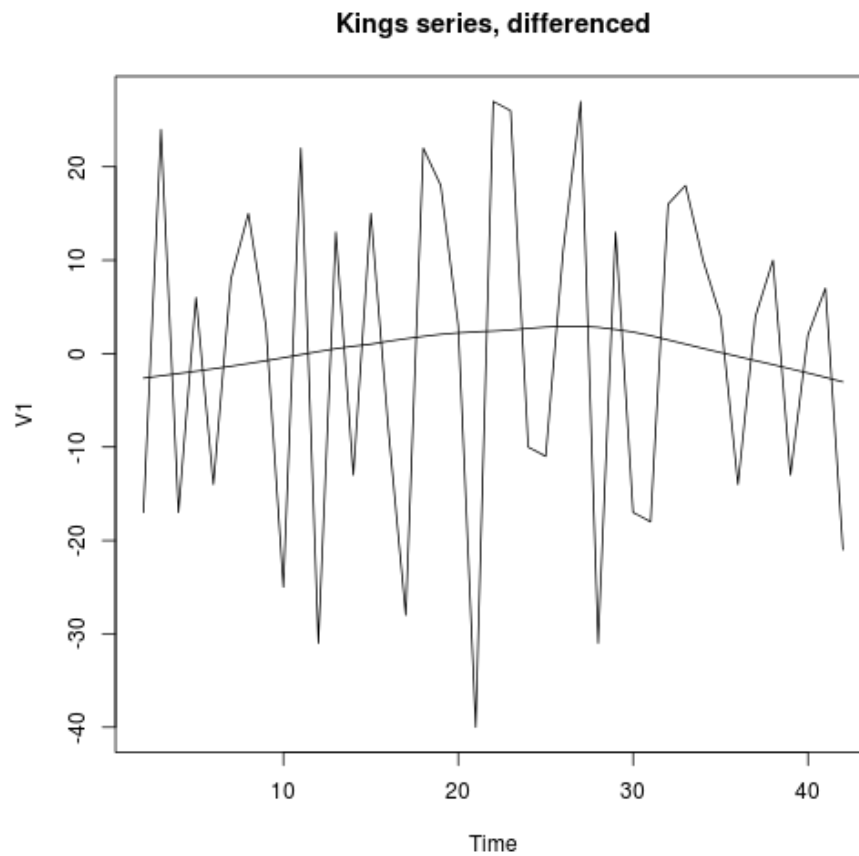


Figure 3:

Births per month in New York City

from January 1946 to December 1959:

```
ny = read.table("nybirths.txt", header = F)
ny.ts = ts(ny, freq = 12, start = c(1946, 1))
```

Note extras on `ts`:

- Time period is 1 year
- 12 observations per year (monthly) in `freq`
- First observation is 1st month of 1946 in `start`

Printing formats nicely

`ny.ts`

```
##           Jan   Feb   Mar   Apr   May   Jun   Jul   Aug   Sep   Oct   Nov
## 1946 26.66 23.60 26.93 24.74 25.81 24.36 24.48 23.90 23.18 23.23 21.67
## 1947 21.44 21.09 23.71 21.67 21.75 20.76 23.48 23.82 23.11 23.11 21.76
## 1948 21.94 20.04 23.59 21.67 22.22 22.12 23.95 23.50 22.24 23.14 21.06
## 1949 21.55 20.00 22.42 20.61 21.76 22.87 24.10 23.75 23.26 22.91 21.52
## 1950 22.60 20.89 24.68 23.67 25.32 23.58 24.67 24.45 24.12 24.25 22.08
## 1951 23.29 23.05 25.08 24.04 24.43 24.67 26.45 25.62 25.01 25.11 22.96
## 1952 23.80 22.27 24.77 22.65 23.99 24.74 26.28 25.82 25.21 25.20 23.16
## 1953 24.36 22.64 25.57 24.06 25.43 24.64 27.01 26.61 26.27 26.46 25.25
## 1954 24.66 23.30 26.98 26.20 27.21 26.12 26.71 26.88 26.15 26.38 24.71
## 1955 24.99 24.24 26.72 23.48 24.77 26.22 28.36 28.60 27.91 27.78 25.69
## 1956 26.22 24.22 27.91 26.98 28.53 27.14 28.98 28.17 28.06 29.14 26.29
## 1957 26.59 24.85 27.54 26.90 28.88 27.39 28.07 28.14 29.05 28.48 26.63
## 1958 27.13 24.92 28.96 26.59 27.93 28.01 29.23 28.76 28.41 27.95 25.91
## 1959 26.08 25.29 27.66 25.95 26.40 25.57 28.86 30.00 29.26 29.01 26.99
##           Dec
## 1946 21.87
## 1947 22.07
## 1948 21.57
## 1949 22.02
## 1950 22.99
## 1951 23.98
## 1952 24.71
## 1953 25.18
## 1954 25.69
```

```
## 1955 26.88
## 1956 26.99
## 1957 27.73
## 1958 26.62
## 1959 27.90
```

Time plot

- Time plot shows extra pattern:

```
plot(ny.ts)
```




Figure 4:

Comments on time plot

- steady increase
- repeating pattern each year (seasonal component).
- Not stationary.

Differencing the New York births

Does differencing help here?

```
ny.diff.ts = diff(ny.ts)
plot(ny.diff.ts)
```

Looks stationary, but some regular spikes.

Decomposing a seasonal time series

Observations for NY births were every month. Are things the same every year?

A visual (using original data):

```
ny.d = decompose(ny.ts)
plot(ny.d)
```

Decomposition bits

Shows:

- original series
- just the trend, going steadily up (except at the start)
- a “seasonal” part: something that repeats every year
- random: what is left over (“residuals”)




Figure 5:




Figure 6:

The seasonal part

Fitted seasonal part is same every year, births lowest in February and highest in July:

```
ny.d$seasonal
```

##		Jan	Feb	Mar	Apr	May	Jun	Jul	Aug
##	1946	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1947	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1948	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1949	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1950	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1951	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1952	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1953	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1954	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1955	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1956	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1957	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1958	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##	1959	-0.6772	-2.0830	0.8625	-0.8017	0.2517	-0.1533	1.4560	1.1646
##		Sep	Oct	Nov	Dec				
##	1946	0.6916	0.7752	-1.1098	-0.3768				
##	1947	0.6916	0.7752	-1.1098	-0.3768				
##	1948	0.6916	0.7752	-1.1098	-0.3768				
##	1949	0.6916	0.7752	-1.1098	-0.3768				
##	1950	0.6916	0.7752	-1.1098	-0.3768				
##	1951	0.6916	0.7752	-1.1098	-0.3768				
##	1952	0.6916	0.7752	-1.1098	-0.3768				
##	1953	0.6916	0.7752	-1.1098	-0.3768				
##	1954	0.6916	0.7752	-1.1098	-0.3768				
##	1955	0.6916	0.7752	-1.1098	-0.3768				
##	1956	0.6916	0.7752	-1.1098	-0.3768				
##	1957	0.6916	0.7752	-1.1098	-0.3768				
##	1958	0.6916	0.7752	-1.1098	-0.3768				
##	1959	0.6916	0.7752	-1.1098	-0.3768				

Time series basics

White noise

Independent random normal. Knowing one value tells you nothing about the next. “Random” process.

```

wn = rnorm(100)
wn.ts = ts(wn)
plot(wn.ts)

```




Figure 7:

Lagging a time series

This means moving a time series one (or more) steps back in time:

```

x = rnorm(5)
x.1 = c(NA, x)
x.0 = c(x, NA)
cbind(x.0, x.1)

```



```
##           x.0      x.1
## [1,] -2.03609      NA
## [2,] -0.57862 -2.03609
## [3,]  0.60836 -0.57862
## [4,]  0.11803  0.60836
## [5,]  0.05634  0.11803
## [6,]          NA  0.05634
```

Have to glue NA onto **both** ends to keep series same length.

Lagging white noise

```
wn.1 = c(NA, wn)
wn.0 = c(wn, NA)
plot(wn.1, wn.0)

cor(wn.1, wn.0, use = "c")

## [1] -0.01676
```

Correlation with lagged series

If you know about white noise at one time point, you know *nothing* about it at the next. This is shown by the scatterplot and the correlation.

On the other hand, this:

```
kings.1 = c(NA, kings.ts)
kings.0 = c(kings.ts, NA)
cor(kings.1, kings.0, use = "c")

## [1] 0.401
```

If one value larger, the next value (a bit) more likely to be larger.

Plot of series vs. lagged series for kings data

```
plot(kings.1, kings.0)
```




Figure 8:




Figure 9:

Two steps back:

```
# one step back
cor(kings.1, kings.0, use = "c")
```

```
## [1] 0.401
```

```
# one step plus one more
kings.2 = c(NA, kings.1)
kings.0 = c(kings.0, NA)
cor(kings.0, kings.2, use = "c")
```

```
## [1] 0.2457
```

Still a correlation two steps back, but smaller.

Autocorrelation

Correlation of time series with *itself* one, two, ... time steps back is useful idea, called **autocorrelation**. Make a plot of it with **acf**.

White noise:

```
acf(wn.ts)
```

No autocorrelations beyond chance, anywhere. Ignore 0.

Autocorrelations work best on *stationary* series.

Kings, differenced

```
acf(kings.diff.ts)
```

Comments on autocorrelations of kings series

Negative autocorrelation at lag 1, nothing beyond that.

- If one value of differenced series positive, next one most likely negative.
- If one king lives longer than predecessor, next one likely lives shorter.




Figure 10:




Figure 11:

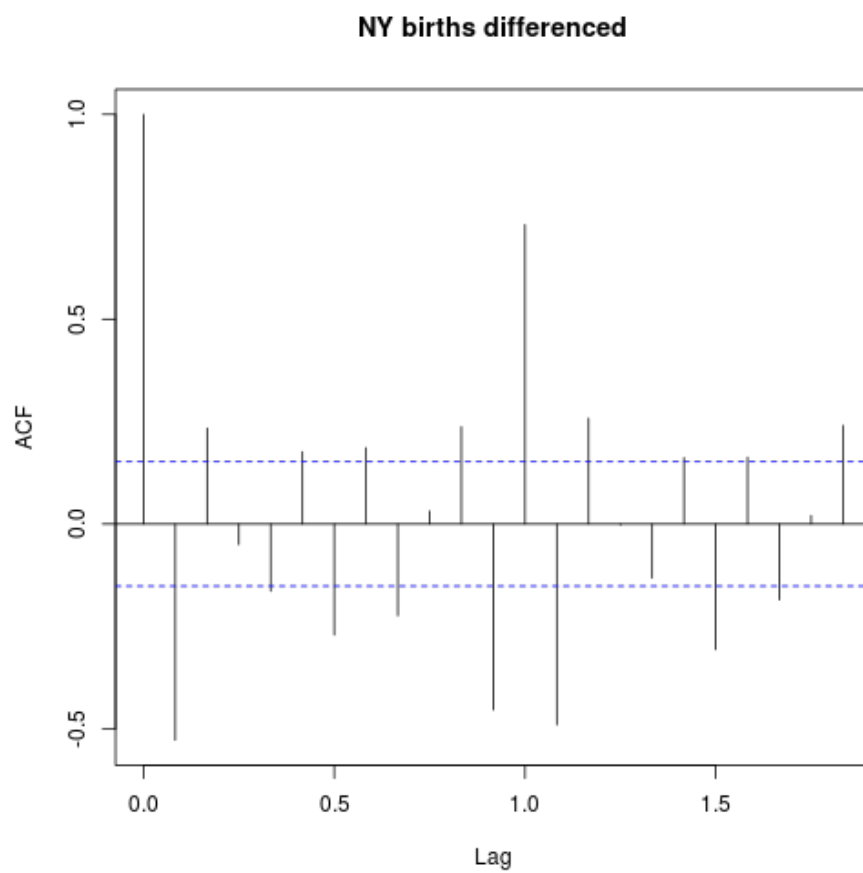


Figure 12:

NY births, differenced

```
acf(ny.diff.ts, main = "NY births differenced")
```

Comments

Lots of stuff:

- large positive autocorrelation at 1.0 years (July one year like July last year)
- large negative autocorrelation at 1 month.
- smallish but significant negative autocorrelation at 0.5 year = 6 months.
- Other stuff – complicated.

Souvenir sales

Monthly sales for a beach souvenir shop in Queensland, Australia:

```
souv = read.table("souvenir.txt", header = F)
souv.ts = ts(souv, frequency = 12, start = 1987)
souv.ts
```

##	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct
## 1987	1665	2398	2841	3547	3753	3715	4350	3566	5022	6423
## 1988	2500	5198	7225	4806	5901	4951	6179	4752	5496	5835
## 1989	4717	5703	9958	5305	6492	6631	7350	8177	8573	9690
## 1990	5921	5815	12421	6370	7609	7225	8121	7979	8093	8477
## 1991	4827	6470	9639	8821	8722	10209	11277	12552	11637	13607
## 1992	7615	9850	14558	11587	9333	13082	16733	19889	23933	25391
## 1993	10243	11267	21827	17357	15998	18602	26155	28587	30505	30821
##	Nov	Dec								
## 1987	7601	19756								
## 1988	12600	28542								
## 1989	15152	34061								
## 1990	17915	30114								
## 1991	21822	45061								
## 1992	36025	80722								
## 1993	46634	104661								




Figure 13:

Plot of souvenir sales

```
plot(souv.ts)
```

Comments

Several problems:

- Mean goes up over time
- Variability gets larger as mean gets larger
- Not stationary

Problem-fixing

Fix non-constant variability first by taking logs:

```
souv.log.ts = log(souv.ts)
plot(souv.log.ts)
```

Differencing

Mean still not constant, so try taking differences:

```
souv.log.diff.ts = diff(souv.log.ts)
plot(souv.log.diff.ts)
```

Comments

- Now stationary
- but clear seasonal effect.

Decomposing to see the seasonal effect

```
souv.d = decompose(souv.log.diff.ts)
plot(souv.d)
```




Figure 14:




Figure 15:




Figure 16:

Comments

Big drop in one month's differences. Look at seasonal component to see which:

```
souv.d$seasonal
```

##	Jan	Feb	Mar	Apr	May	Jun	Jul
## 1987		0.23293	0.49069	-0.39701	0.02410	0.05074	0.13553
## 1988	-1.90372	0.23293	0.49069	-0.39701	0.02410	0.05074	0.13553
## 1989	-1.90372	0.23293	0.49069	-0.39701	0.02410	0.05074	0.13553
## 1990	-1.90372	0.23293	0.49069	-0.39701	0.02410	0.05074	0.13553
## 1991	-1.90372	0.23293	0.49069	-0.39701	0.02410	0.05074	0.13553
## 1992	-1.90372	0.23293	0.49069	-0.39701	0.02410	0.05074	0.13553
## 1993	-1.90372	0.23293	0.49069	-0.39701	0.02410	0.05074	0.13553
##	Aug	Sep	Oct	Nov	Dec		
## 1987	-0.03710	0.08651	0.09148	0.47311	0.75274		
## 1988	-0.03710	0.08651	0.09148	0.47311	0.75274		
## 1989	-0.03710	0.08651	0.09148	0.47311	0.75274		
## 1990	-0.03710	0.08651	0.09148	0.47311	0.75274		
## 1991	-0.03710	0.08651	0.09148	0.47311	0.75274		
## 1992	-0.03710	0.08651	0.09148	0.47311	0.75274		
## 1993	-0.03710	0.08651	0.09148	0.47311	0.75274		

Autocorrelations

```
acf(souv.log.diff.ts)
```

- Big positive autocorrelation at 1 year (strong seasonal effect)
- Small negative autocorrelation at 1 and 2 months.

Moving average

A particular type of time series called a **moving average** or MA process captures idea of autocorrelations at a few lags but not at others.

Here's generation of MA(1) process, with autocorrelation at lag 1 but not otherwise:

```
e = rnorm(100)
y = numeric(0)
y[1] = 0
beta = 1
```




Figure 17:

```

for (i in 2:100) {
  y[i] = e[i] + beta * e[i - 1]
}
y

##   [1]  0.000000  1.459947  1.003636  0.290535  0.927762  0.698584  1.081925
##   [8]  1.199785 -0.755392 -2.501267 -1.491053  0.316963  0.007674  0.915412
##  [15]  0.699441 -0.709542 -1.510910 -1.931463 -2.699014 -0.191624  2.068134
##  [22]  1.953255  1.102764  1.160782  0.354067 -2.057886 -2.156346 -1.848551
##  [29] -0.683571 -1.433488 -2.226428 -1.569811 -1.998215 -0.774296 -0.776534
##  [36] -1.896003  0.593475  1.769262 -0.002125  0.350905  0.301290  0.257268
##  [43]  0.814916  0.174229 -1.457311 -0.440137  0.937195  0.621751  1.770630
##  [50]  1.167974 -1.342307 -1.565790 -0.992549  0.868109  2.128187  0.196433
##  [57]  0.345667  1.989314  0.467650 -0.891247  1.781103  0.104384 -1.702983
##  [64] -0.912250 -0.756331 -0.902107 -2.832956 -2.894871 -1.981636 -1.571337
##  [71] -0.648132 -1.068589 -2.818508 -2.154688  0.154307 -1.036771 -0.407083
##  [78] -0.419829 -1.736900 -0.042858  1.651501 -0.099512 -0.327476  0.890043
##  [85] -0.913421 -1.161211 -0.860284  0.261762  1.662999  1.536472  0.373325
##  [92]  1.756376  2.459355 -0.055969  0.467825  1.075871 -2.575363 -1.257956
##  [99]  0.583707 -0.714033

```

Comments

- `e` contains independent “random shocks”.
- Start process at 0.
- Then, each value of the time series has that time’s random shock, plus a multiple of the last time’s random shock.
- `y[i]` has shock in common with `y[i-1]`; should be a lag 1 autocorrelation.
- But `y[i]` has no shock in common with `y[i-2]`, so no lag 2 autocorrelation (or beyond).

ACF for MA(1) process

```
acf(y)
```

As promised, everything beyond lag 1 is just chance.

AR process

Another kind of time series is AR process, where each value depends on previous one, like this:




Figure 18:

```

e = rnorm(100)
x = numeric(0)
x[1] = 0
alpha = 0.7
for (i in 2:100) {
  x[i] = alpha * x[i - 1] + e[i]
}
x

##   [1]  0.00000  0.69150 -0.27157 -1.69374 -0.04625 -0.61290  0.26465
##   [8] -0.21494 -1.31429  0.44277  0.09918  0.19081 -1.02379  0.16694
##  [15]  0.98375  0.04866  1.22332 -0.04785 -0.21368 -0.68229  0.25079
##  [22] -0.86025  1.75818  1.19266  0.30513  2.41225  1.28151  1.68979
##  [29]  2.01816  3.53755  1.85841  2.32514  1.77112  2.12224  0.91096
##  [36]  1.58477  2.08225  1.09623 -0.76369 -0.70810 -1.84440 -0.38985
##  [43] -1.04266 -0.86988 -1.14486 -3.18900 -2.93376 -2.16076 -1.59509
##  [50] -1.74905 -3.13933 -3.02637 -1.44219 -1.55490 -1.73929 -2.00996
##  [57] -2.66272 -3.20338 -3.51822 -3.07147 -3.97834 -3.76372 -3.52533
##  [64] -3.45189 -0.06075 -0.57178  0.81558 -0.27386  0.75055 -1.41071
##  [71] -2.60771 -0.77008 -0.44599  0.92660 -0.50866 -0.28001 -0.69942
##  [78] -0.87488 -1.34524 -1.24758 -2.20687 -1.55319 -0.03080 -0.30484
##  [85]  1.32564  1.13382  0.88142  0.19973 -1.03974 -0.60656  0.27269
##  [92]  0.49555  0.74140  0.41685 -0.01248 -0.08956  1.09794  0.51406
##  [99]  1.27608  0.05862

```

Comments

- Each random shock now only used for its own value of x
- but $x[i]$ also depends on previous value $x[i-1]$
- so correlated with previous value
- *but* $x[i]$ also contains multiple of $x[i-2]$ and previous x 's
- so all x 's correlated, but autocorrelation dying away.

ACF for AR(1) series

```
acf(x)
```




Figure 19:

Partial autocorrelation function

This cuts off for an AR series:

`pacf(x)`




Figure 20:

The lag-2 autocorrelation should not be significant, and is only just.

PACF for an MA series

Decays slowly for an MA series:

`pacf(y)`




Figure 21:

The old way of doing time series analysis

Starting from a series with constant variability (eg. transform first to get it, as for souvenirs):

- Assess stationarity.
- If not stationary, take differences as many times as needed until it is.
- Look at ACF, see if it dies off. If it does, you have MA series.
- Look at PACF, see if that dies off. If it does, have AR series.
- If neither dies off, probably have a mixed “ARMA” series.
- Fit coefficients (like regression slopes).
- Do forecasts.

The new way of doing time series analysis (in R)

- Transform series if needed to get constant variability
- Use package `forecast`.
- Use function `auto.arima` to estimate what kind of series best fits data.
- Use `forecast` to see what will happen in future.

Anatomy of `auto.arima` output

```
library(forecast)

## This is forecast 4.03

auto.arima(y)

## Warning: p-value greater than printed p-value

## Series: y
## ARIMA(0,0,1) with zero mean
##
## Coefficients:
##          ma1
##          0.907
## s.e.    0.062
##
## sigma^2 estimated as 0.978:  log likelihood=-141.6
## AIC=287.3   AICc=287.4   BIC=292.5
```

- ARIMA part tells you what kind of series you are estimated to have:
- first number (first 0) is AR (autoregressive) part
- second number (second 0) is amount of differencing here
- third number (1) is MA (moving average) part
- Below that, coefficients (with SEs)
- AICc is measure of fit (lower better)

What other models were possible?

Run `auto.arima` with `trace=T`:

```
auto.arima(y, trace = T)

## Warning: p-value greater than printed p-value

##
## ARIMA(2,0,2) with non-zero mean : 1e+20
## ARIMA(0,0,0) with non-zero mean : 345.2
## ARIMA(1,0,0) with non-zero mean : 314
## ARIMA(0,0,1) with non-zero mean : 287.9
## ARIMA(1,0,1) with non-zero mean : 290.1
## ARIMA(0,0,2) with non-zero mean : 290.1
## ARIMA(1,0,2) with non-zero mean : 291.8
## ARIMA(0,0,1) with zero mean      : 287.4
## ARIMA(1,0,1) with zero mean      : 289.5
## ARIMA(0,0,0) with zero mean      : 346.1
## ARIMA(0,0,2) with zero mean      : 289.5
## ARIMA(1,0,2) with zero mean      : 290.6
##
## Best model: ARIMA(0,0,1) with zero mean

## Series: y
## ARIMA(0,0,1) with zero mean
##
## Coefficients:
##          ma1
##          0.907
## s.e.    0.062
##
## sigma^2 estimated as 0.978:  log likelihood=-141.6
## AIC=287.3   AICc=287.4   BIC=292.5
```

Also possible were MA(2) and ARMA(1,1), both with AICc=289.5.

Doing it all the new way

White noise

```
wn.aa = auto.arima(wn.ts)

## Warning: p-value greater than printed p-value

wn.aa

## Series: wn.ts
## ARIMA(0,0,0) with zero mean
##
## sigma^2 estimated as 1.11: log likelihood=-147.2
## AIC=296.3   AICc=296.4   BIC=298.9
```

Best fit *is* white noise (no AR, no MA, no differencing).

Forecasts

```
forecast(wn.aa)

##      Point Forecast   Lo 80 Hi 80   Lo 95 Hi 95
## 101              0 -1.351 1.351 -2.066 2.066
## 102              0 -1.351 1.351 -2.066 2.066
## 103              0 -1.351 1.351 -2.066 2.066
## 104              0 -1.351 1.351 -2.066 2.066
## 105              0 -1.351 1.351 -2.066 2.066
## 106              0 -1.351 1.351 -2.066 2.066
## 107              0 -1.351 1.351 -2.066 2.066
## 108              0 -1.351 1.351 -2.066 2.066
## 109              0 -1.351 1.351 -2.066 2.066
## 110              0 -1.351 1.351 -2.066 2.066
```

Forecasts all 0, since the past doesn't help to predict future.

MA(1)

```
y.aa = auto.arima(y)

## Warning: p-value greater than printed p-value
```



```

y.aa

## Series: y
## ARIMA(0,0,1) with zero mean
##
## Coefficients:
##          ma1
##          0.907
## s.e.    0.062
##
## sigma^2 estimated as 0.978:  log likelihood=-141.6
## AIC=287.3   AICc=287.4   BIC=292.5

y.f = forecast(y.aa)

```

Plotting the forecasts for MA(1)

```
plot(y.f)
```

AR(1)

```

x.aa = auto.arima(x)

## Warning: p-value greater than printed p-value

## Warning: possible convergence problem: optim gave code=1

## Warning: possible convergence problem: optim gave code=1

x.aa

## Series: x
## ARIMA(0,1,1)
##
## Coefficients:
##          ma1
##          -0.354
## s.e.    0.106
##
## sigma^2 estimated as 0.969:  log likelihood=-139
## AIC=282   AICc=282.1   BIC=287.2

```




Figure 22:

Oops!

Got it wrong! Fit right AR(1) model:

```
x.arima = arima(x, order = c(1, 0, 0))
x.arima

## Series: x
## ARIMA(1,0,0) with non-zero mean
##
## Coefficients:
##          ar1  intercept
##          0.776    -0.365
## s.e.    0.061      0.422
##
## sigma^2 estimated as 0.957:  log likelihood=-140.2
## AIC=286.3   AICc=286.6   BIC=294.1
```

Forecasts for x

```
plot(forecast(x.arima))
```

Comparing wrong model

```
plot(forecast(x.aa))
```

Kings

```
kings.aa = auto.arima(kings.ts)

## Warning: p-value greater than printed p-value

kings.aa

## Series: kings.ts
## ARIMA(0,1,1)
##
## Coefficients:
##          ma1
##          -0.722
## s.e.    0.121
##
## sigma^2 estimated as 230:  log likelihood=-170.1
## AIC=344.1   AICc=344.4   BIC=347.6
```

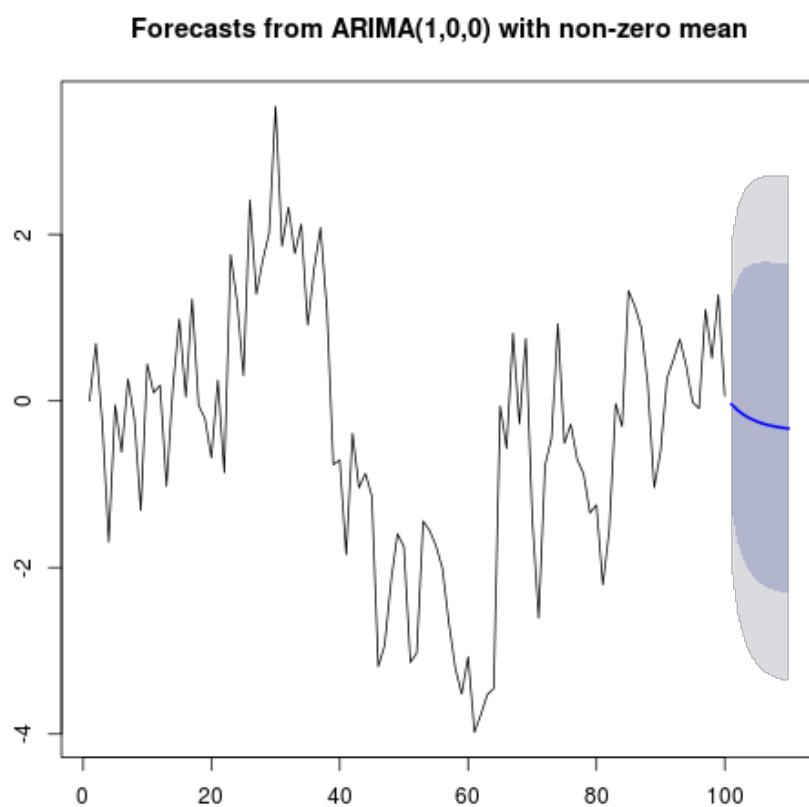


Figure 23:




Figure 24:

Kings forecasts

```
kings.f = forecast(kings.aa)
kings.f
```

```
##      Point Forecast Lo 80 Hi 80 Lo 95 Hi 95
## 43          67.75 48.30 87.20 38.00 97.50
## 44          67.75 47.56 87.94 36.87 98.63
## 45          67.75 46.84 88.66 35.78 99.72
## 46          67.75 46.16 89.35 34.72 100.78
## 47          67.75 45.49 90.01 33.70 101.80
## 48          67.75 44.84 90.66 32.71 102.79
## 49          67.75 44.21 91.29 31.75 103.76
## 50          67.75 43.59 91.91 30.81 104.70
## 51          67.75 42.99 92.51 29.89 105.61
## 52          67.75 42.41 93.09 29.00 106.51
```

Kings forecasts, plotted

```
plot(kings.f)
```

NY births

```
ny.aa = auto.arima(ny.ts)
```

```
## Warning: p-value greater than printed p-value
```

```
ny.aa
```

```
## Series: ny.ts
## ARIMA(2,1,2)(1,1,1)[12]
##
## Coefficients:
##          ar1      ar2      ma1      ma2      sar1      sma1
##          0.654  -0.454  -0.726   0.253  -0.243  -0.845
## s.e.    0.300    0.243    0.323   0.288   0.099   0.099
##
## sigma^2 estimated as 0.392:  log likelihood=-157.4
## AIC=328.9   AICc=329.7   BIC=350.2
```

```
ny.f = forecast(ny.aa, h = 36)
```

Going 36 time periods (3 years) into future.




Figure 25:

NY births forecasts

Not *quite* same every year:

ny.f

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 1960	27.69	26.89	28.49	26.46	28.92
## Feb 1960	26.08	24.98	27.17	24.40	27.75
## Mar 1960	29.27	28.04	30.49	27.39	31.14
## Apr 1960	27.59	26.29	28.90	25.60	29.59
## May 1960	28.93	27.55	30.32	26.82	31.05
## Jun 1960	28.55	27.07	30.03	26.29	30.82
## Jul 1960	29.85	28.27	31.43	27.43	32.27
## Aug 1960	29.45	27.78	31.13	26.89	32.01
## Sep 1960	29.16	27.41	30.92	26.48	31.85
## Oct 1960	29.21	27.38	31.05	26.41	32.01
## Nov 1960	27.26	25.36	29.17	24.35	30.18
## Dec 1960	28.07	26.09	30.05	25.04	31.09
## Jan 1961	27.67	25.64	29.70	24.56	30.78
## Feb 1961	26.21	24.13	28.30	23.02	29.40
## Mar 1961	29.23	27.09	31.37	25.95	32.50
## Apr 1961	27.58	25.38	29.78	24.22	30.94
## May 1961	28.71	26.46	30.96	25.27	32.15
## Jun 1961	28.22	25.92	30.52	24.70	31.74
## Jul 1961	29.99	27.64	32.34	26.39	33.58
## Aug 1961	29.96	27.56	32.36	26.29	33.63
## Sep 1961	29.57	27.12	32.01	25.82	33.31
## Oct 1961	29.55	27.05	32.04	25.73	33.36
## Nov 1961	27.58	25.04	30.12	23.69	31.47
## Dec 1961	28.41	25.82	31.00	24.45	32.37
## Jan 1962	28.05	25.39	30.72	23.98	32.13
## Feb 1962	26.56	23.83	29.29	22.38	30.74
## Mar 1962	29.62	26.82	32.41	25.34	33.89
## Apr 1962	27.96	25.11	30.81	23.60	32.32
## May 1962	29.15	26.24	32.05	24.70	33.59
## Jun 1962	28.68	25.72	31.64	24.15	33.21
## Jul 1962	30.33	27.31	33.35	25.71	34.95
## Aug 1962	30.22	27.14	33.29	25.51	34.92
## Sep 1962	29.85	26.72	32.98	25.06	34.63
## Oct 1962	29.85	26.66	33.03	24.98	34.71
## Nov 1962	27.88	24.65	31.12	22.94	32.83
## Dec 1962	28.71	25.42	31.99	23.68	33.73

Plotting the forecasts

```
plot(ny.f)
```




Figure 26:

Log-souvenir sales

```
souv.aa = auto.arima(souv.log.ts)
souv.aa

## Series: souv.log.ts
## ARIMA(2,0,0)(1,1,0)[12] with drift
##
```

```
## Coefficients:
##          ar1      ar2      sar1  drift
##          0.349  0.360 -0.328  0.025
## s.e.      0.109  0.116   0.133  0.004
##
## sigma^2 estimated as 0.03:  log likelihood=23.04
## AIC=-36.09   AICc=-35.18   BIC=-24.71

souv.f = forecast(souv.aa, h = 27)
```

The forecasts

Differenced series showed low value for January (large drop). December highest, Jan and Feb lowest:

```
souv.f
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 1994	9.561	9.339	9.783	9.221	9.901
## Feb 1994	9.710	9.475	9.946	9.351	10.070
## Mar 1994	10.274	10.015	10.532	9.878	10.669
## Apr 1994	10.042	9.775	10.308	9.634	10.450
## May 1994	9.912	9.638	10.185	9.493	10.330
## Jun 1994	10.121	9.844	10.398	9.697	10.545
## Jul 1994	10.428	10.149	10.708	10.000	10.856
## Aug 1994	10.543	10.262	10.824	10.113	10.973
## Sep 1994	10.646	10.363	10.928	10.214	11.078
## Oct 1994	10.671	10.388	10.954	10.238	11.104
## Nov 1994	11.063	10.780	11.346	10.629	11.497
## Dec 1994	11.870	11.586	12.154	11.436	12.304
## Jan 1995	9.850	9.525	10.175	9.353	10.348
## Feb 1995	9.981	9.650	10.312	9.475	10.487
## Mar 1995	10.576	10.236	10.916	10.056	11.096
## Apr 1995	10.345	10.002	10.688	9.820	10.870
## May 1995	10.230	9.884	10.577	9.701	10.760
## Jun 1995	10.420	10.073	10.768	9.888	10.953
## Jul 1995	10.739	10.390	11.088	10.205	11.272
## Aug 1995	10.845	10.495	11.194	10.310	11.380
## Sep 1995	10.935	10.585	11.285	10.399	11.470
## Oct 1995	10.955	10.605	11.306	10.419	11.491
## Nov 1995	11.354	11.004	11.705	10.818	11.891
## Dec 1995	12.162	11.811	12.513	11.625	12.698
## Jan 1996	10.149	9.755	10.544	9.546	10.753
## Feb 1996	10.286	9.886	10.686	9.675	10.898
## Mar 1996	10.871	10.461	11.280	10.244	11.497

Plotting the forecasts

```
plot(souv.f)
```

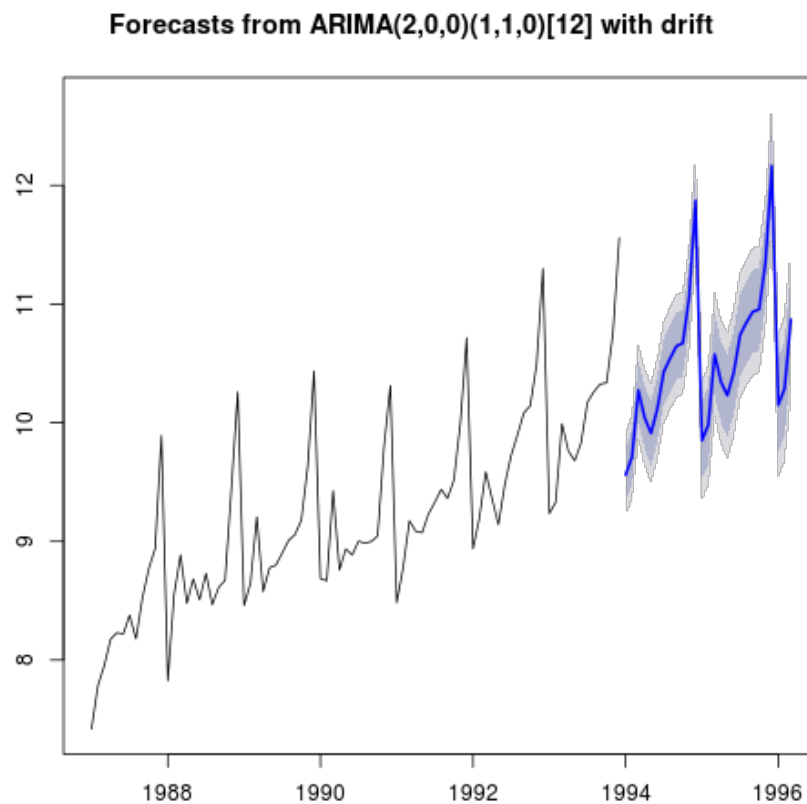


Figure 27:

Global mean temperatures, revisited

```
attach(temp)
temp.ts = ts(temperature, start = 1880)
temp.aa = auto.arima(temp.ts)
```

```
## Warning: p-value smaller than printed p-value
```

```
## Warning: p-value greater than printed p-value

temp.aa

## Series: temp.ts
## ARIMA(1,1,1) with drift
##
## Coefficients:
##          ar1      ma1  drift
##      0.260  -0.783  0.007
## s.e.  0.116   0.069  0.003
##
## sigma^2 estimated as 0.00902:  log likelihood=121.3
## AIC=-234.6   AICc=-234.2   BIC=-223.1
```

Forecasts

```
temp.f = forecast(temp.aa)
plot(temp.f)
```

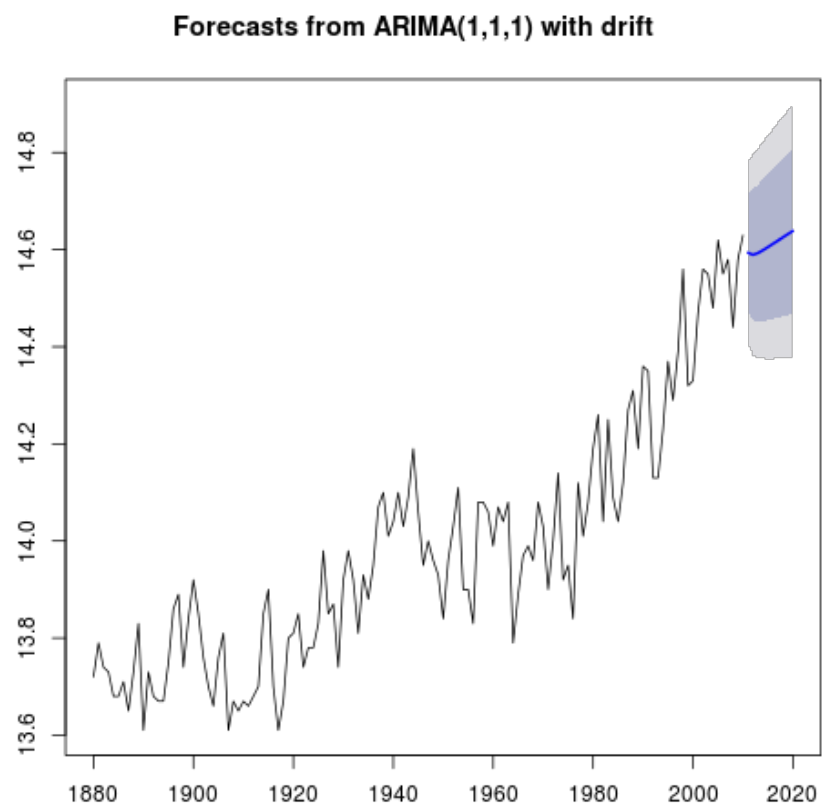


Figure 28: